

**Note to other teachers and users of these slides:** We would be delighted if you found our material useful for giving your own lectures. Feel free to use these slides verbatim, or to modify them to fit your own needs. If you make use of a significant portion of these slides in your own lecture, please include this message, or a link to our web site: <http://cs224w.Stanford.edu>

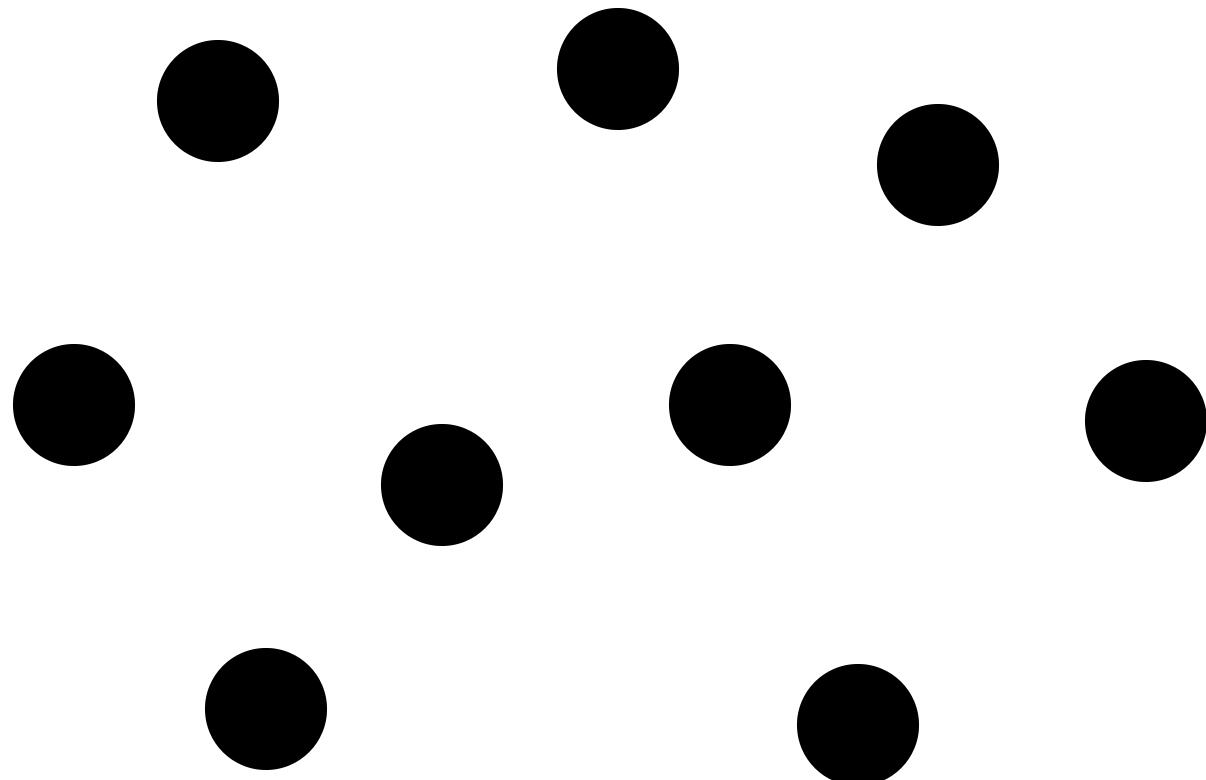
# Stanford CS224W: Machine Learning with Graphs Winter 2022/23

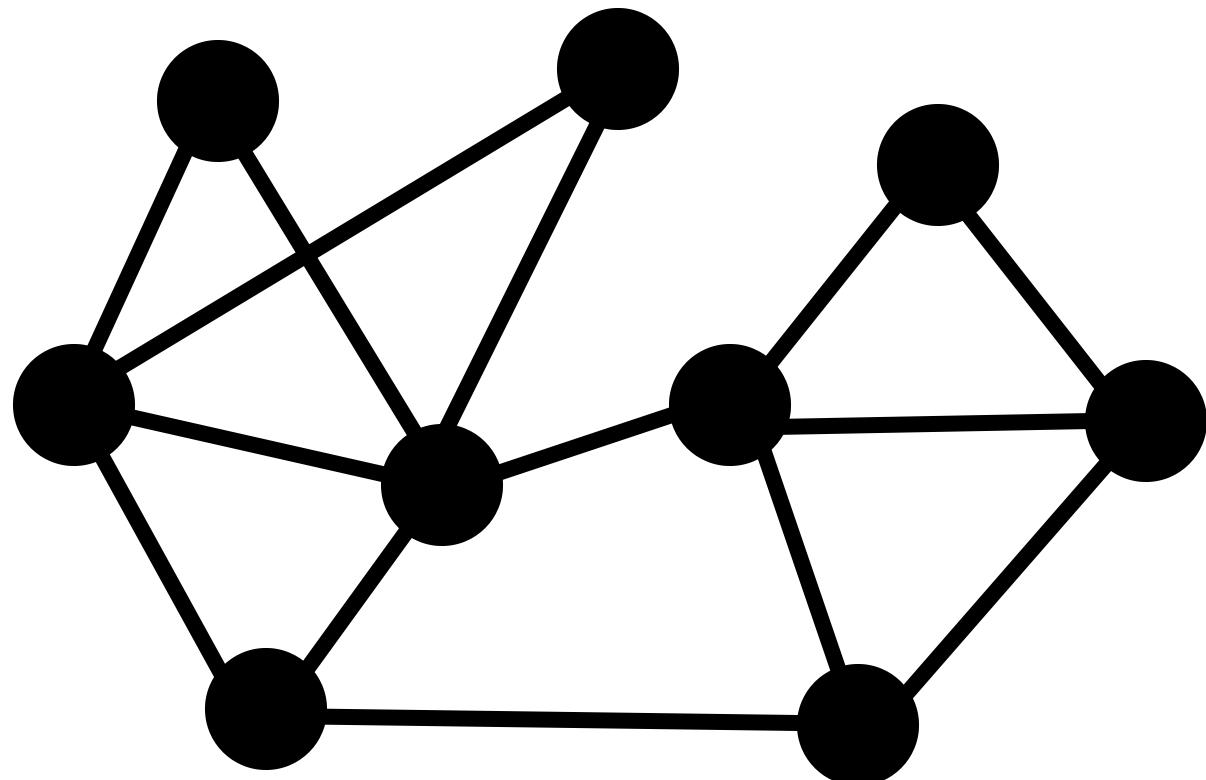
CS224W: Machine Learning with Graphs  
Jure Leskovec, Stanford University  
<http://cs224w.stanford.edu>



# Why Graphs?

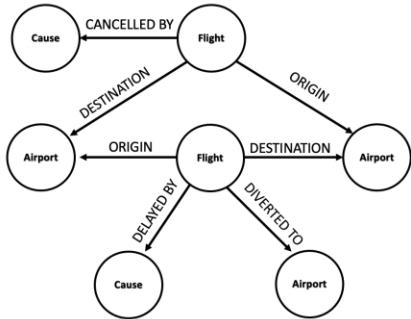
Graphs are a general language for describing and analyzing entities with relations/interactions





# Graph

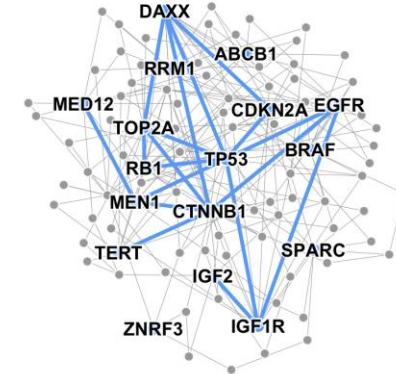
# Many Types of Data are Graphs (1)



## Event Graphs



## Computer Networks



## Disease Pathways

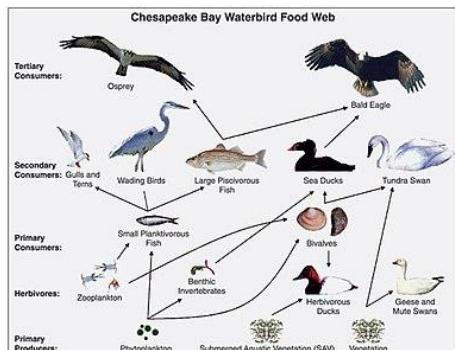


Image credit: [Wikipedia](#)

## Food Webs



Image credit: [Pinterest](#)

## Particle Networks



Image credit: [visitlondon.com](#)

## Underground Networks

# Many Types of Data are Graphs (2)



Image credit: [Medium](#)

**Social Networks**

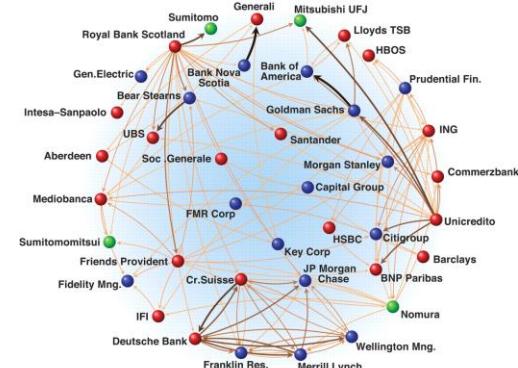


Image credit: [Science](#)



Image credit: [Lumen Learning](#)

**Economic Networks**    **Communication Networks**

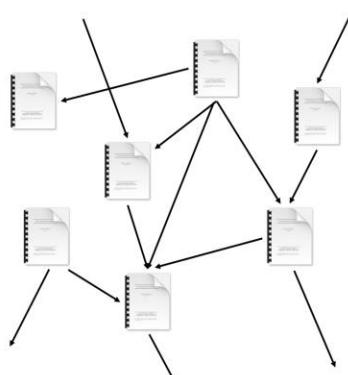


Image credit: [Missoula Current News](#)

**Citation Networks**

**Internet**

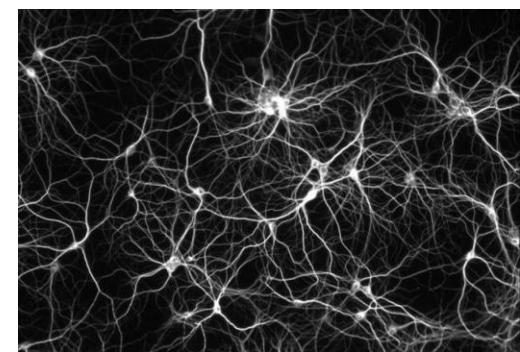


Image credit: [The Conversation](#)

**Networks of Neurons**

# Many Types of Data are Graphs (3)

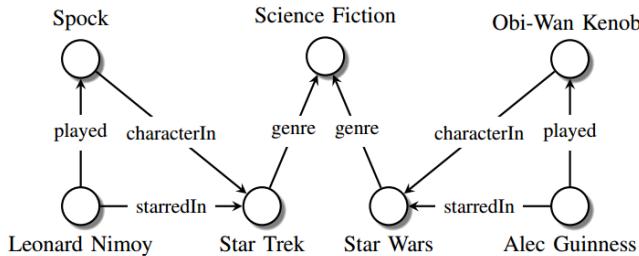


Image credit: Maximilian Nickel et al

## Knowledge Graphs

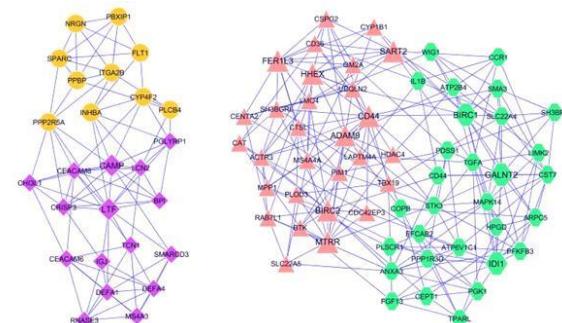


Image credit: ese.wustl.edu

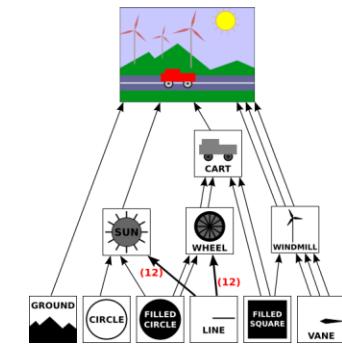


Image credit: math.hws.edu

## Scene Graphs

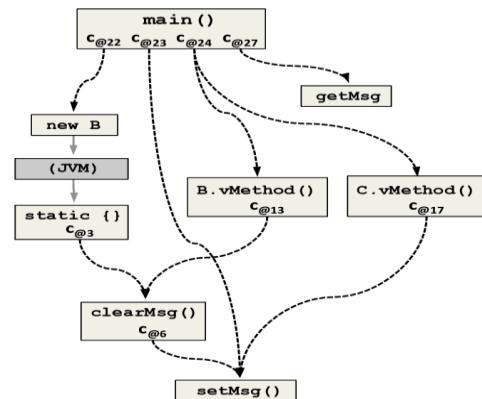


Image credit: ResearchGate

## Code Graphs

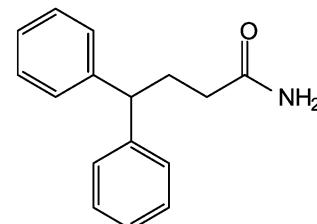


Image credit: MDPI

## Molecules

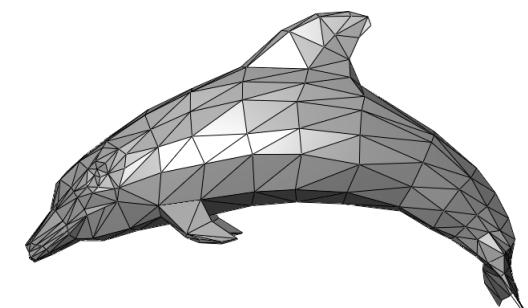
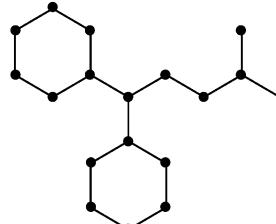


Image credit: Wikipedia

## 3D Shapes

# Graphs and Relational Data

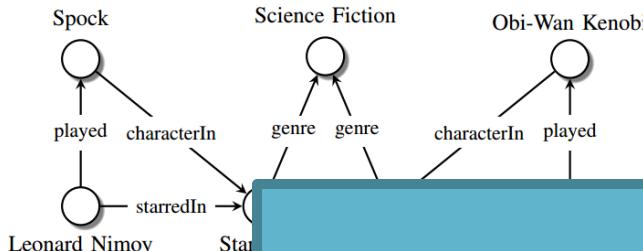


Image credit:

Knowledge

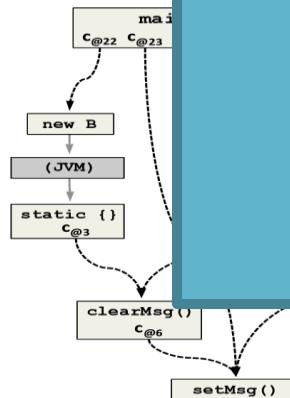


Image credit: ResearchGate

Code Graphs

Main question:

How do we take advantage of relational structure for better prediction?

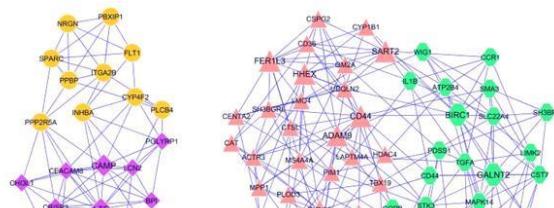
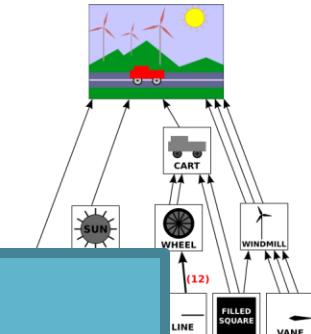


Image credit: MDPI

Molecules



math.hws.edu

Graphs

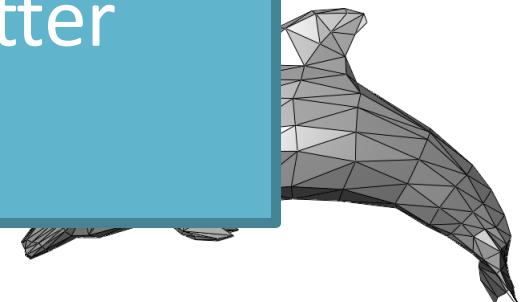


Image credit: Wikipedia

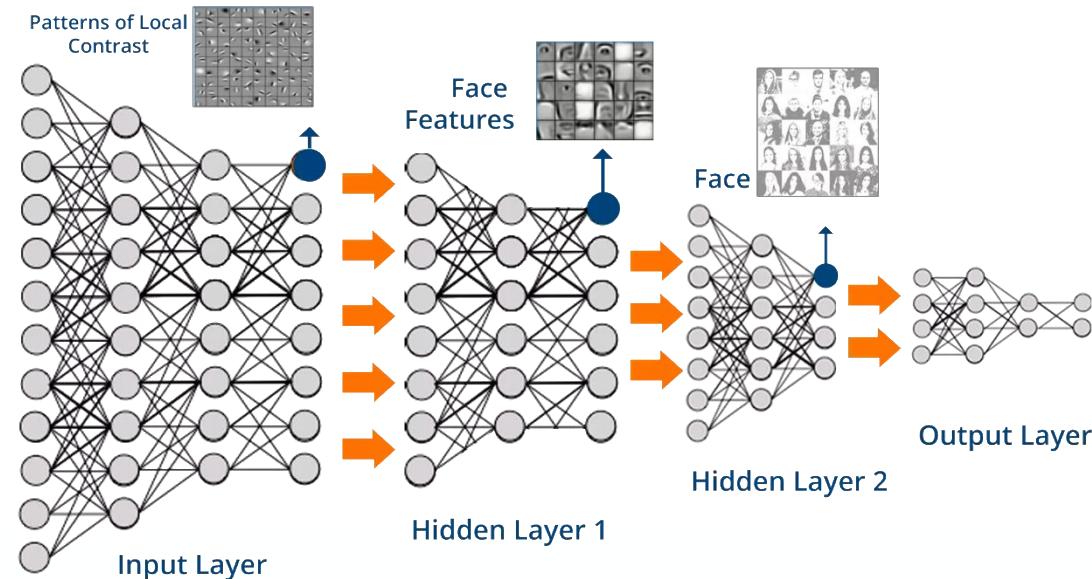
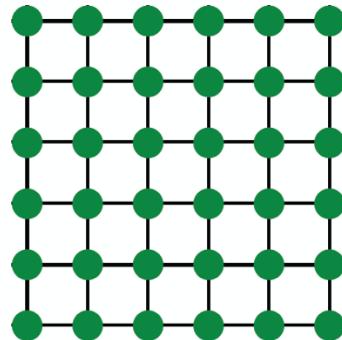
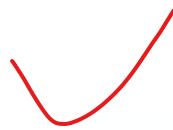
3D Shapes

# Graphs: Machine Learning

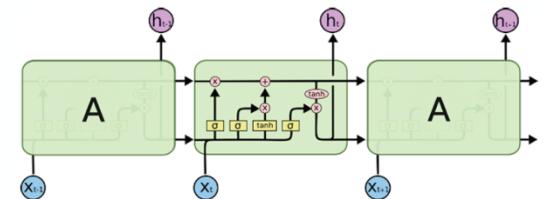
Complex domains have a rich relational structure, which can be represented as a **relational graph**

**By explicitly modeling relationships we achieve better performance!**

# Today: Modern ML Toolbox



**Text/Speech**



Modern deep learning toolbox is designed  
for simple sequences & grids

Doubt thou the stars are fire,  
Doubt that the sun doth move;  
Doubt truth to be a liar;  
But never doubt I love...

Text



Audio signals



Images

Modern  
deep learning toolbox  
is designed for  
sequences & grids

Not everything  
can be represented as  
a sequence or a grid

**How can we develop neural  
networks that are much more  
broadly applicable?**

New frontiers beyond classic neural  
networks that only learn on images  
and sequences

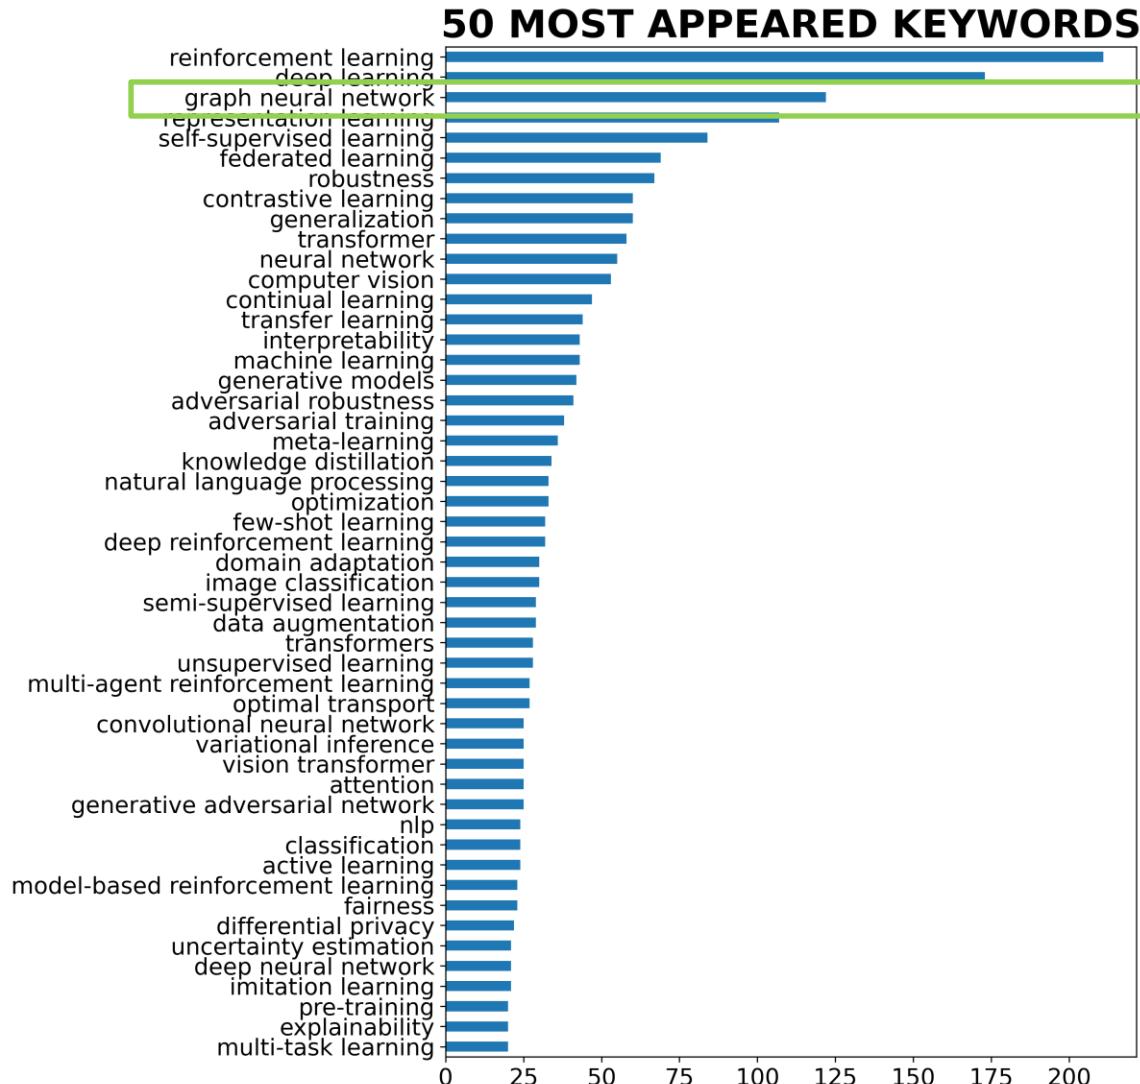
# This Class

Graphs are the new frontier  
of deep learning

Graphs connect things.

# Hot subfield in ML

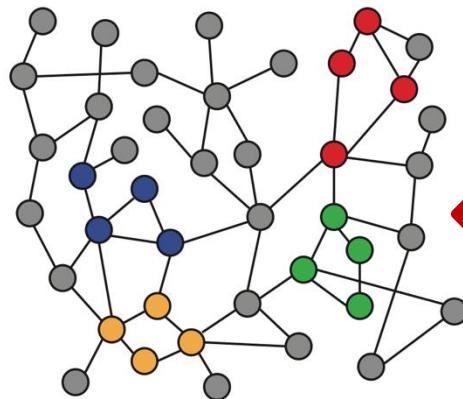
ICLR 2022 keywords



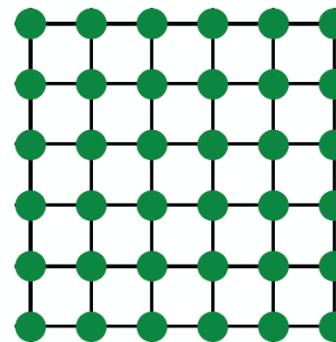
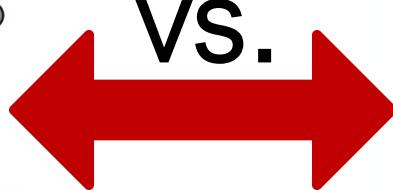
# Why is Graph Deep Learning Hard?

✓ Networks are complex.

- Arbitrary size and complex topological structure (*i.e.*, no spatial locality like grids)



Networks



Images



Text

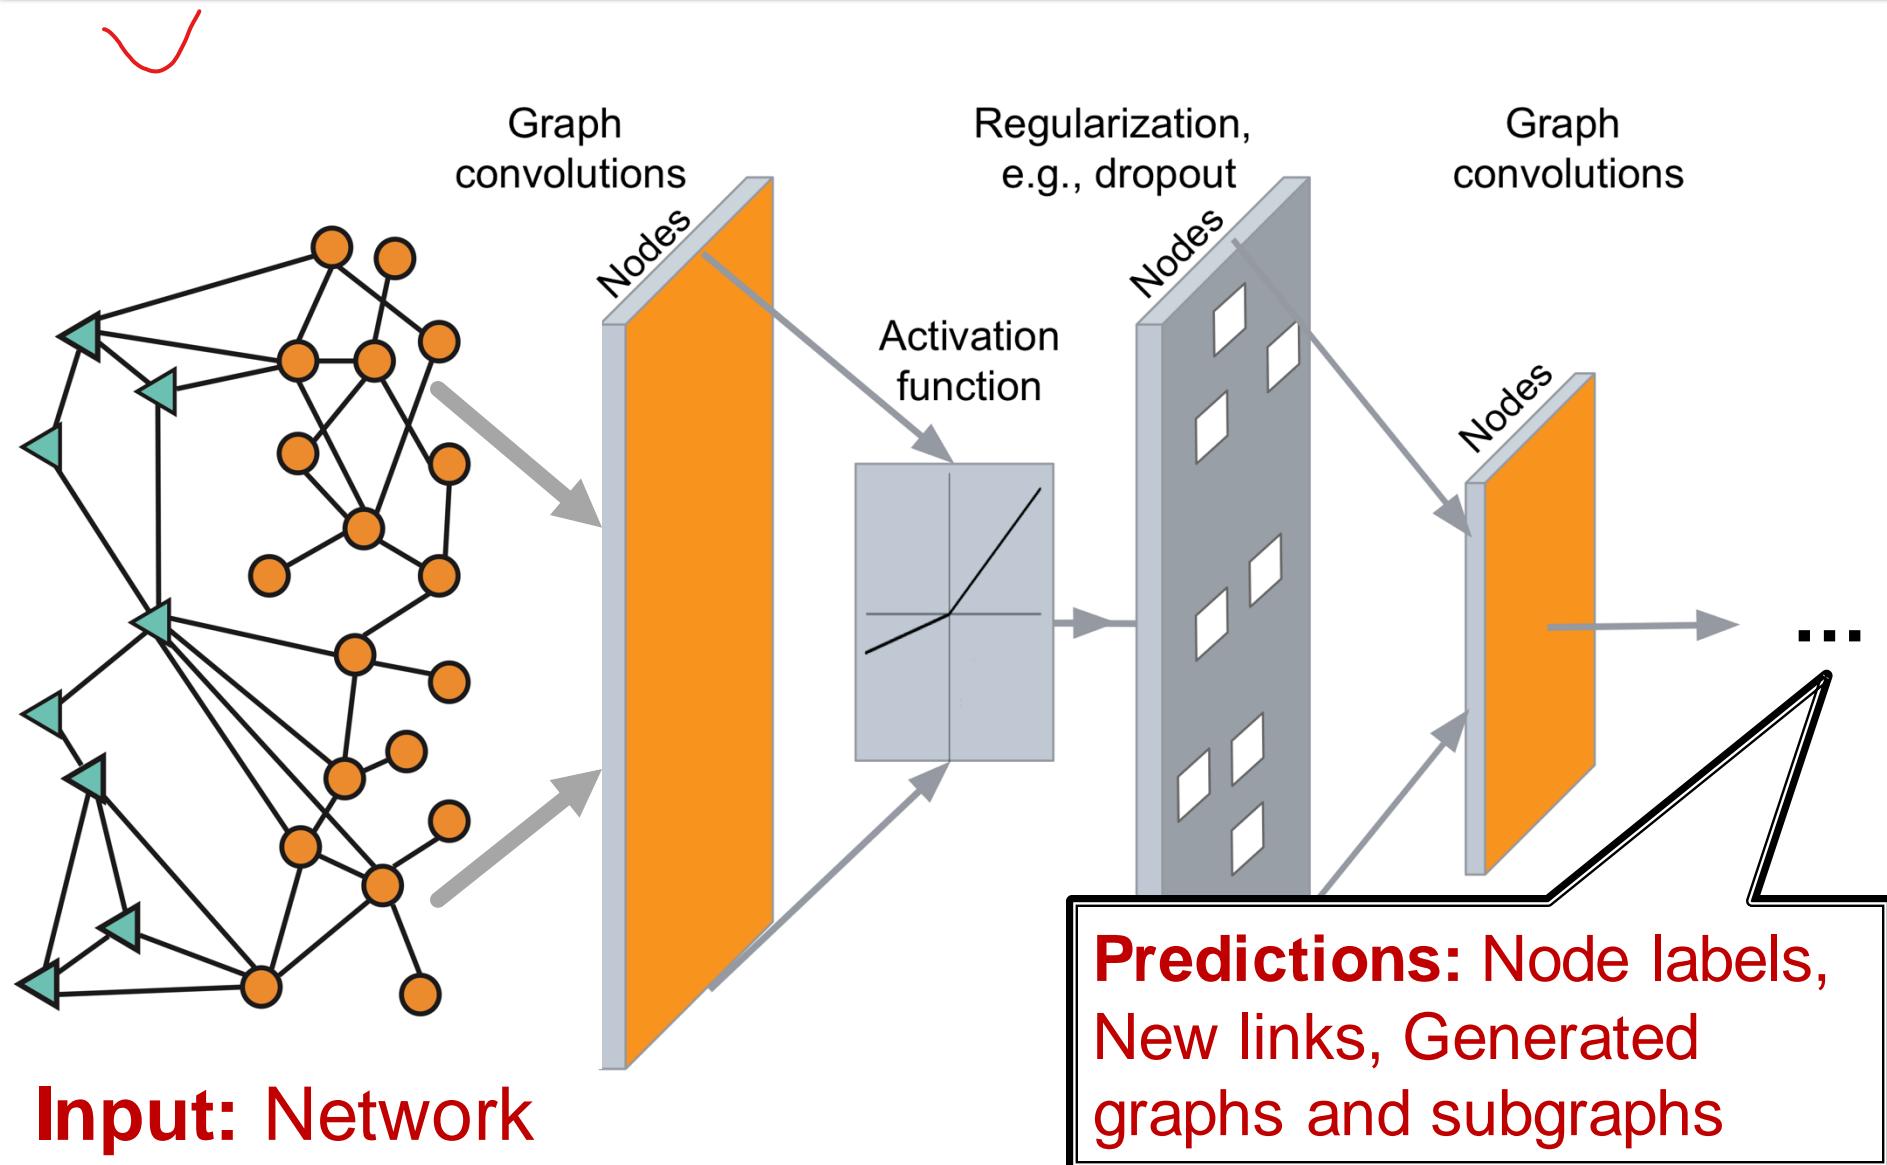
- No fixed node ordering or reference point
- Often dynamic and have multimodal features

# This Course

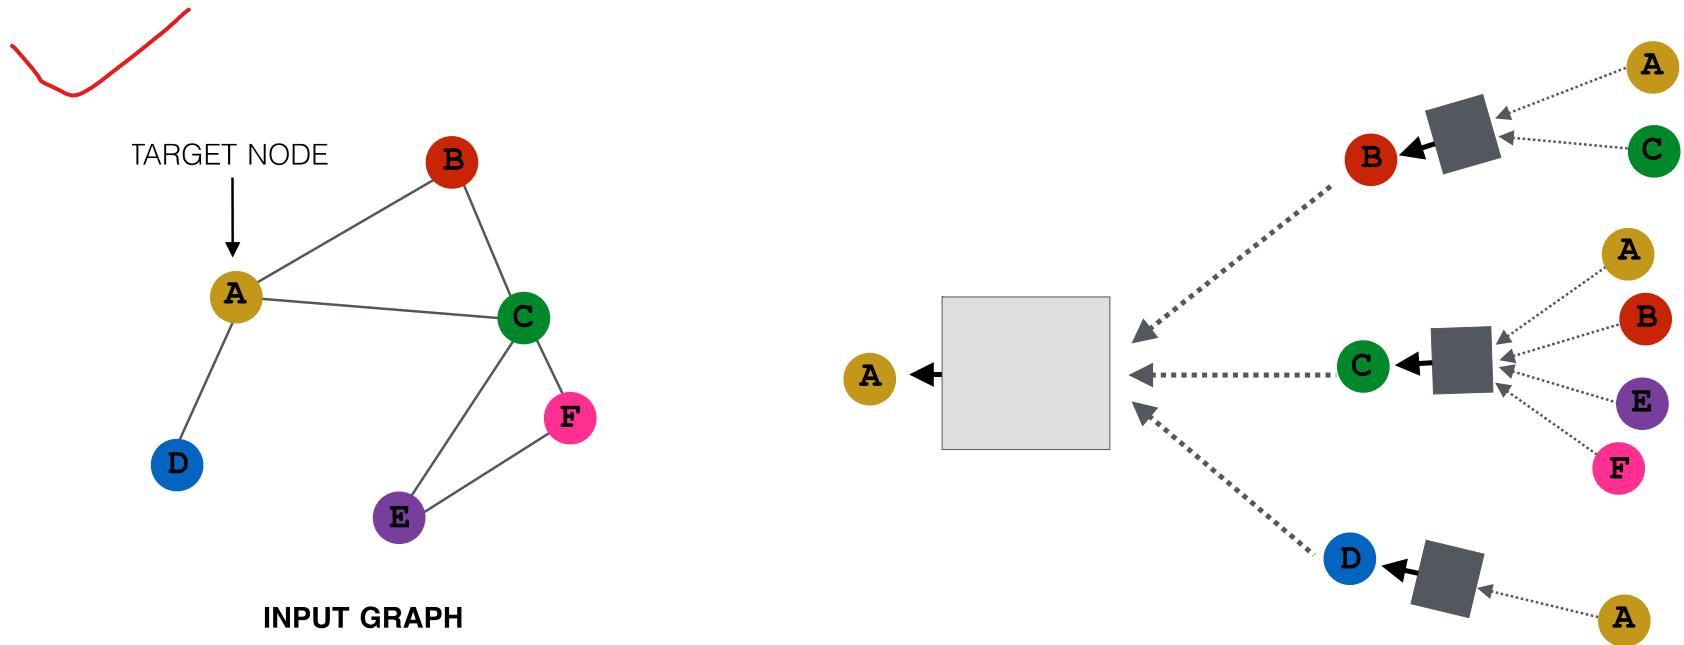
How can we develop neural networks  
that are much more broadly  
applicable?

Graphs are the new frontier  
of deep learning

# CS224W: ML withGraphs



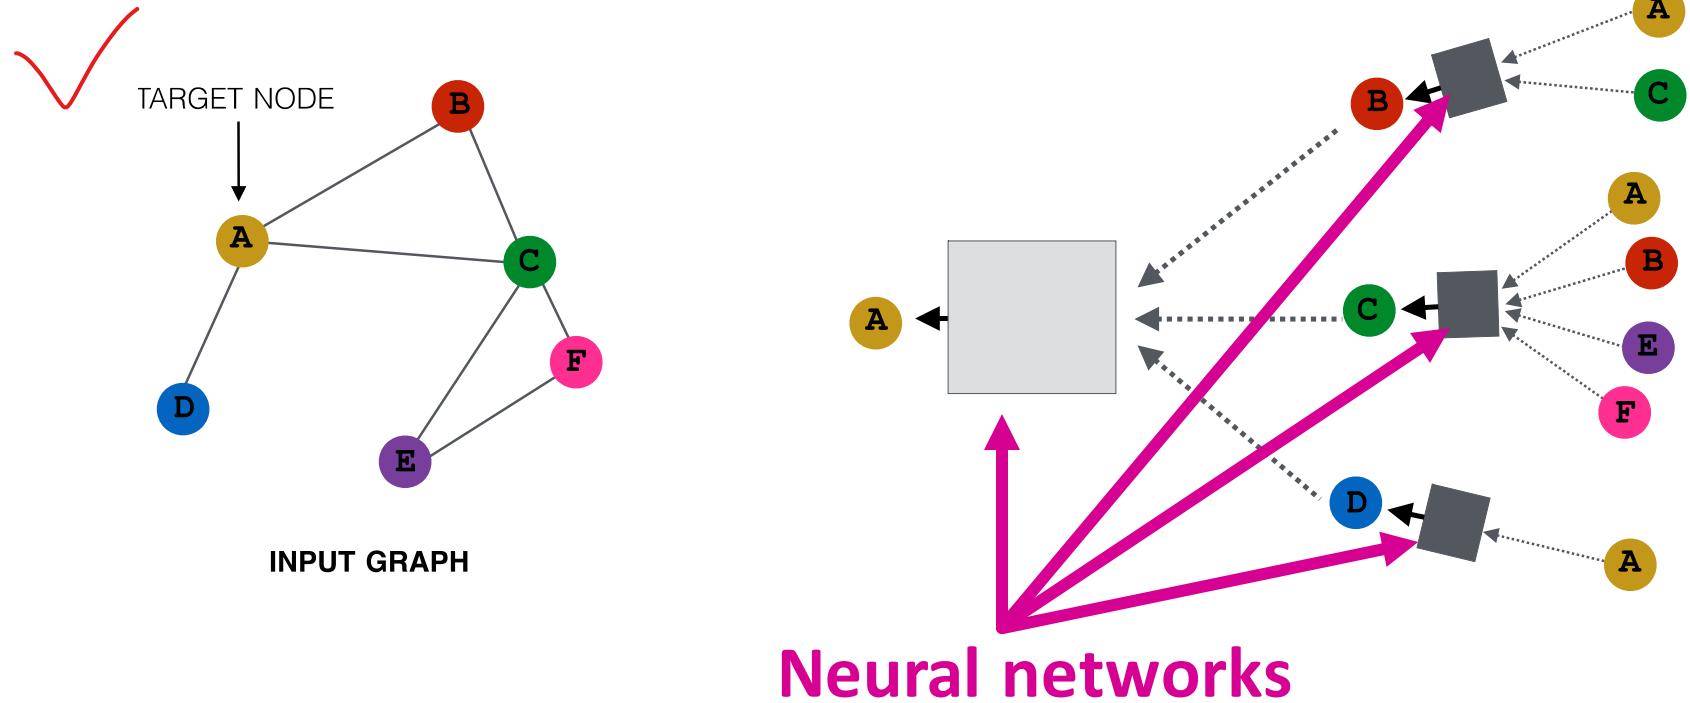
# Graph Neural Networks



Each node defines a computation graph

- Each edge in this graph is a transformation/aggregation function

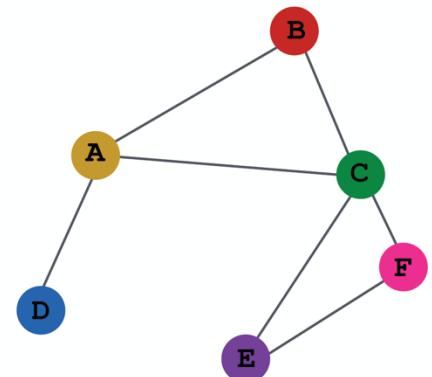
# Graph Neural Networks



**Intuition:** Nodes aggregate information from their neighbors using neural networks

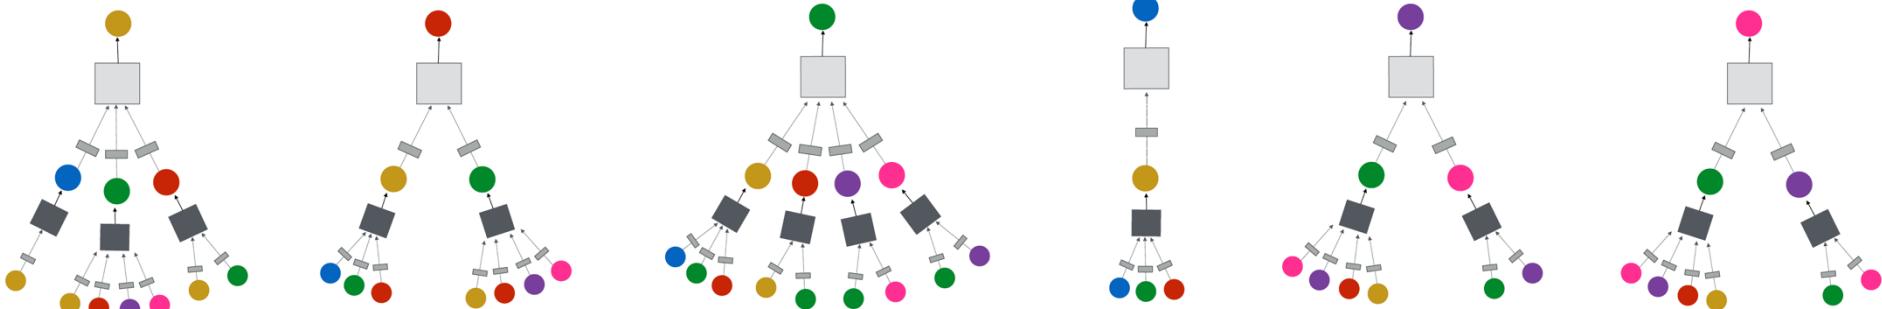
# Idea: Aggregate Neighbors

✓ **Intuition:** Network neighborhood defines a computation graph



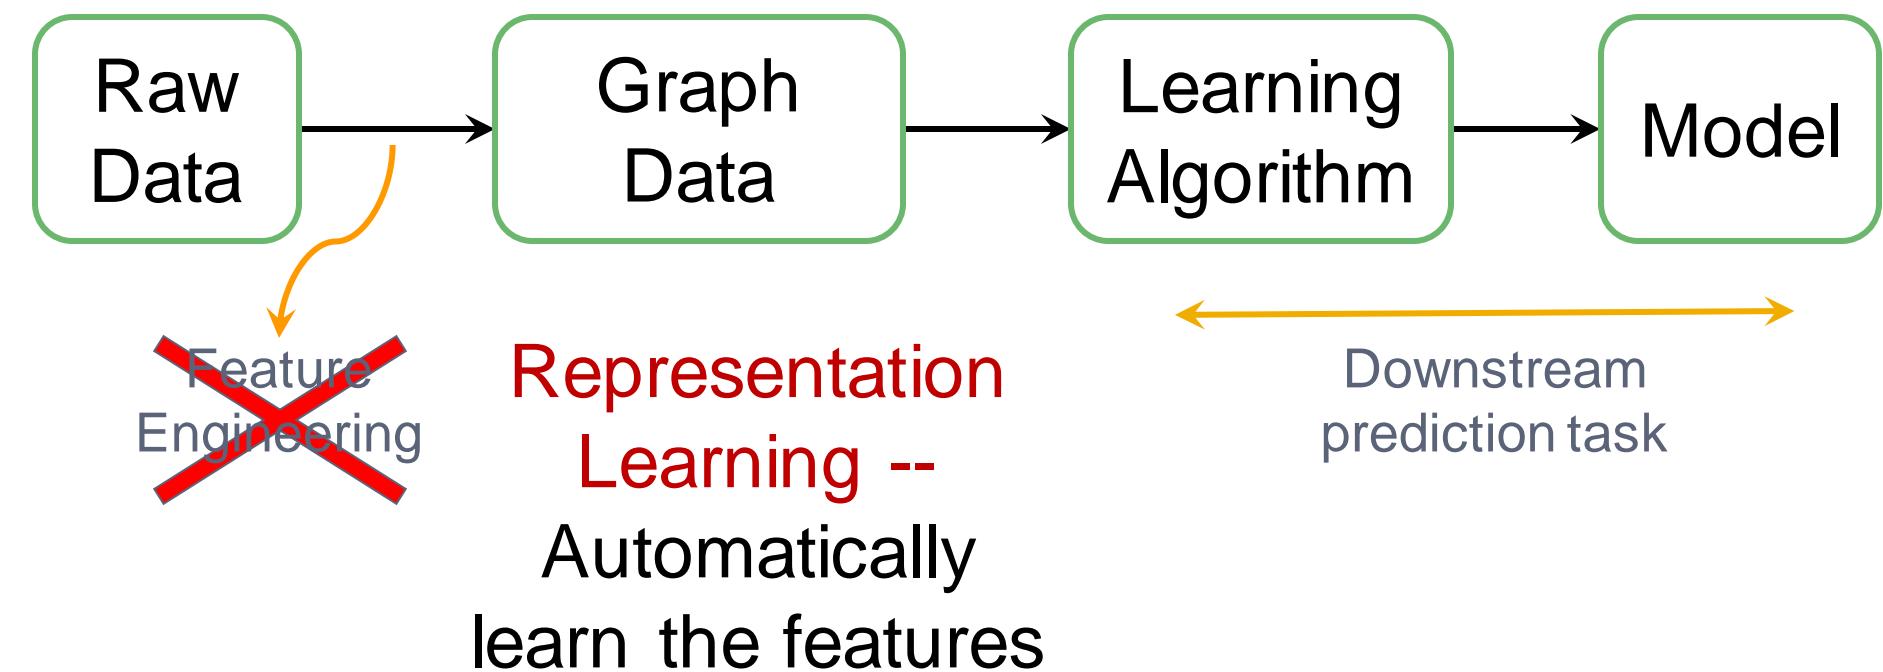
INPUT GRAPH

Every node defines a computation graph based on its neighborhood!



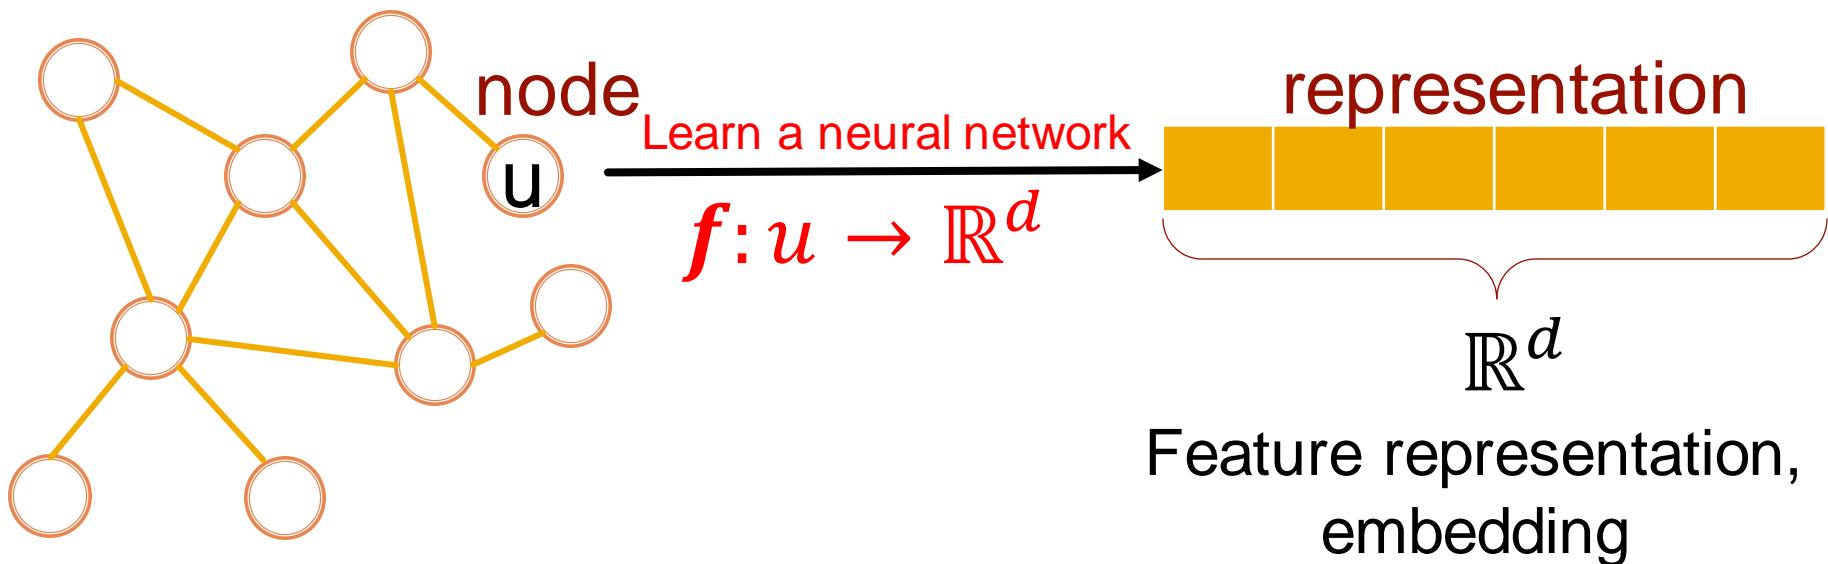
# CS224W & Representation Learning

(Supervised) Machine Learning Lifecycle:  
This feature, that feature. **Every single time!**



# CS224W & Representation Learning

Map nodes to d-dimensional embeddings such that similar nodes in the network are embedded close together



# CS224W Course Outline

We are going to explore Machine Learning and Representation Learning for graph data:

- Traditional methods: Graphlets, Graph Kernels
- Methods for node embeddings: DeepWalk, Node2Vec
- Graph Neural Networks: GCN, GraphSAGE, GAT, Theory of GNNs
- Knowledge graphs and reasoning: TransE, BetaE
- Deep generative models for graphs: GraphRNN
- Applications to Biomedicine, Science, Technology

# CS224W Course Outline

Date	Topic	Date	Topic
Tue, 1/10	1. Introduction; Machine Learning for Graphs	Tue, 2/14	11. Community Structure in Networks
Thu, 1/12	2. Node Embeddings	Thu, 2/16	12. Traditional Generative Models for Graphs
Tue, 1/17	3. Label Propagation for Node Classification	Tue, 2/21	13. Deep Generative Models for Graphs
Thu, 1/19	4. Graph Neural Networks 1: GNN Model	Thu, 2/23	14. Advanced Topics on GNNs
Tue, 1/24	5. Graph Neural Networks 2: Design Space	Tue, 2/28	15. Scaling up GNNs
Thu, 1/26	6. Applications of Graph Neural Networks	Thu, 3/2	16. Explainability
Tue, 1/31	7. Theory of Graph Neural Networks	Tue, 3/7	EXAM
Thu, 2/2	8. Knowledge Graph Embeddings	Thu, 3/9	17. Guest lecture: TBD
Tue, 2/7	9. Reasoning over Knowledge Graphs	Tue, 3/14	18. GNNs for Science
Thu, 2/9	10. Frequent Subgraph Mining with GNNs	Thu, 3/16	19. Special topics in GNNs

# **Stanford CS224W: Applications of Graph ML**

CS224W: Machine Learning with Graphs

Jure Leskovec, Stanford University

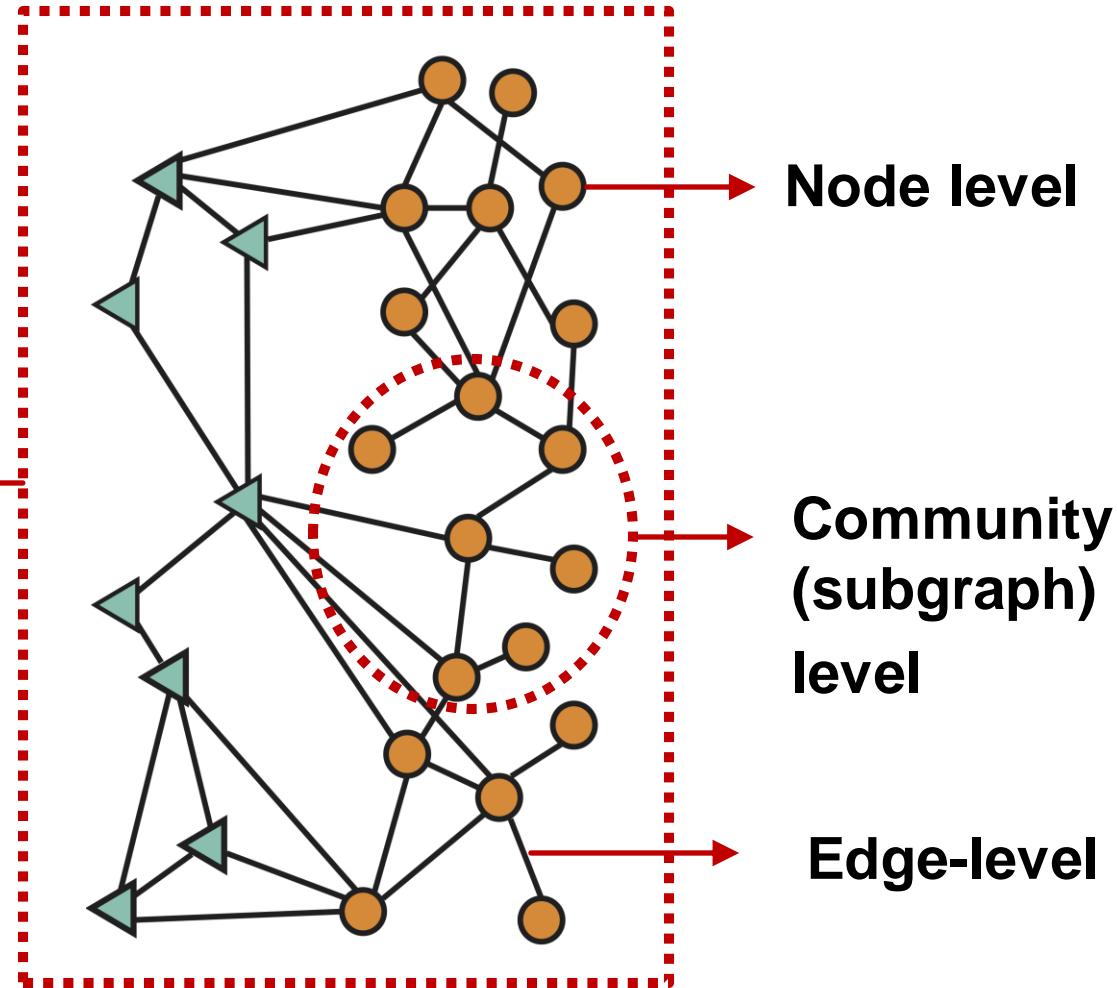
<http://cs224w.stanford.edu>



# Different Types of Tasks

✓

Graph-level  
prediction,  
Graph  
generation



# Classic Graph ML Tasks

- ✓
  - **Node classification**: Predict a property of a node
    - Example: Categorize online users / items
  - **Link prediction**: Predict whether there are missing links between two nodes
    - Example: Knowledge graph completion
  - **Graph classification**: Categorize different graphs
    - Example: Molecule property prediction
  - **Clustering**: Detect if nodes form a community
    - Example: Social circle detection
  - **Other tasks**:
    - **Graph generation**: Drug discovery
    - **Graph evolution**: Physical simulation

# Classic Graph ML Tasks

- **Node classification:** Predict a property of a node
    - Example: Categorize online users / items
  - **Link prediction:** Predict whether there are missing links
    - Example: Predict if two users are friends
  - **Graph classification:** Predict a property of a graph
    - Example: Predict if a graph is social or not
  - **Clustering:** Group nodes into clusters
    - Example: Group users by interests
  - **Others:**
    - **Graph generation:** Drug discovery
    - **Graph evolution:** Physical simulation
- These Graph ML tasks lead to high-impact applications!

# Example of Node-level ML Tasks

# Example (1): Protein Folding

A protein chain acquires its native 3D structure

Every protein is made up of a sequence of amino acids bonded together

These amino acids interact locally to form shapes like helices and sheets

These shapes fold up on larger scales to form the full three-dimensional protein structure

Proteins can interact with other proteins, performing functions such as signalling and transcribing DNA

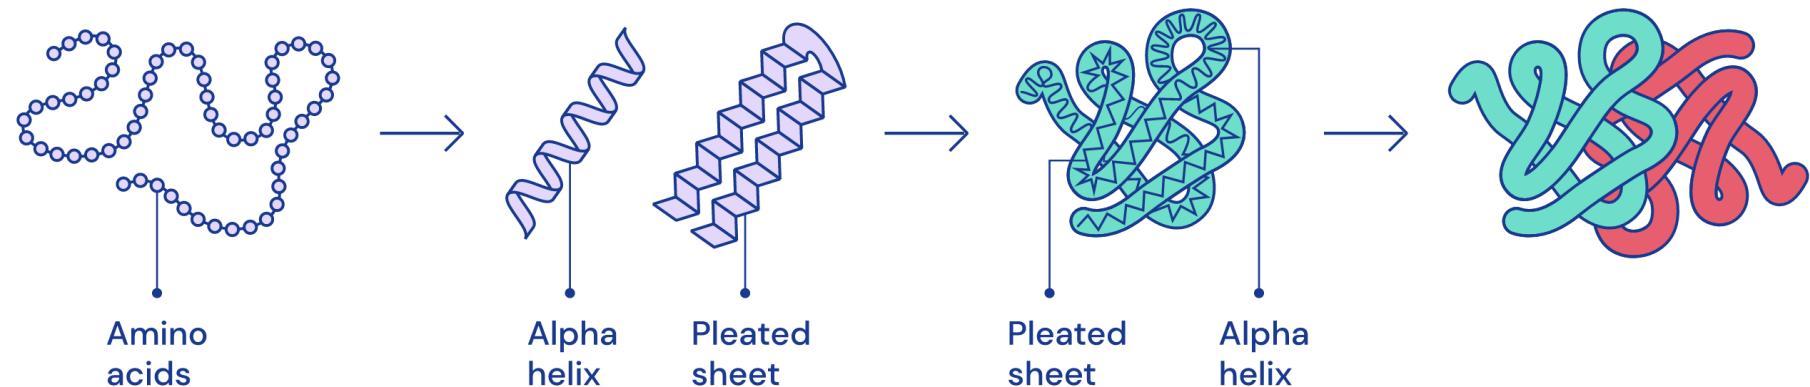


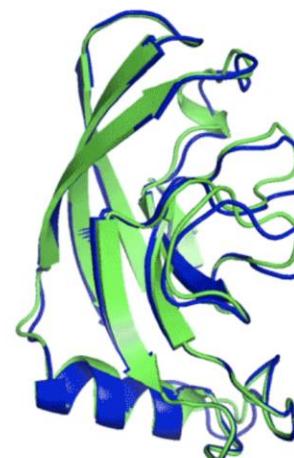
Image credit: [DeepMind](#)

# The Protein Folding Problem

Computationally predict a protein's 3D structure based solely on its amino acid sequence



T1037 / 6vr4  
90.7 GDT  
(RNA polymerase domain)



T1049 / 6y4f  
93.3 GDT  
(adhesin tip)

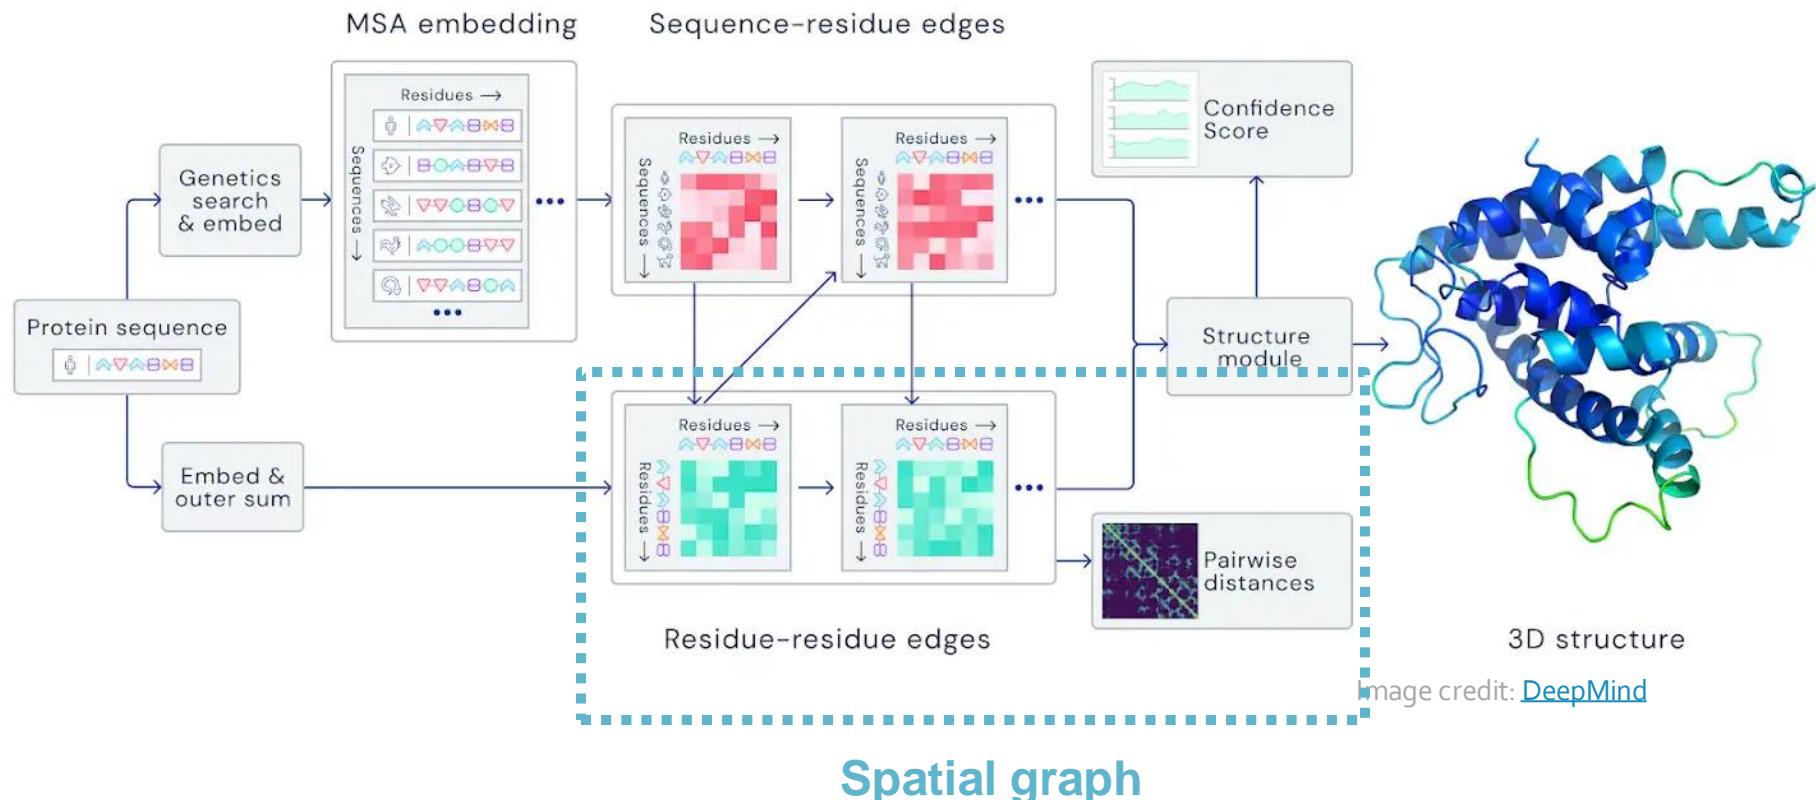
- Experimental result
- Computational prediction

Image credit: [DeepMind](#)



# AlphaFold: Solving Protein Folding

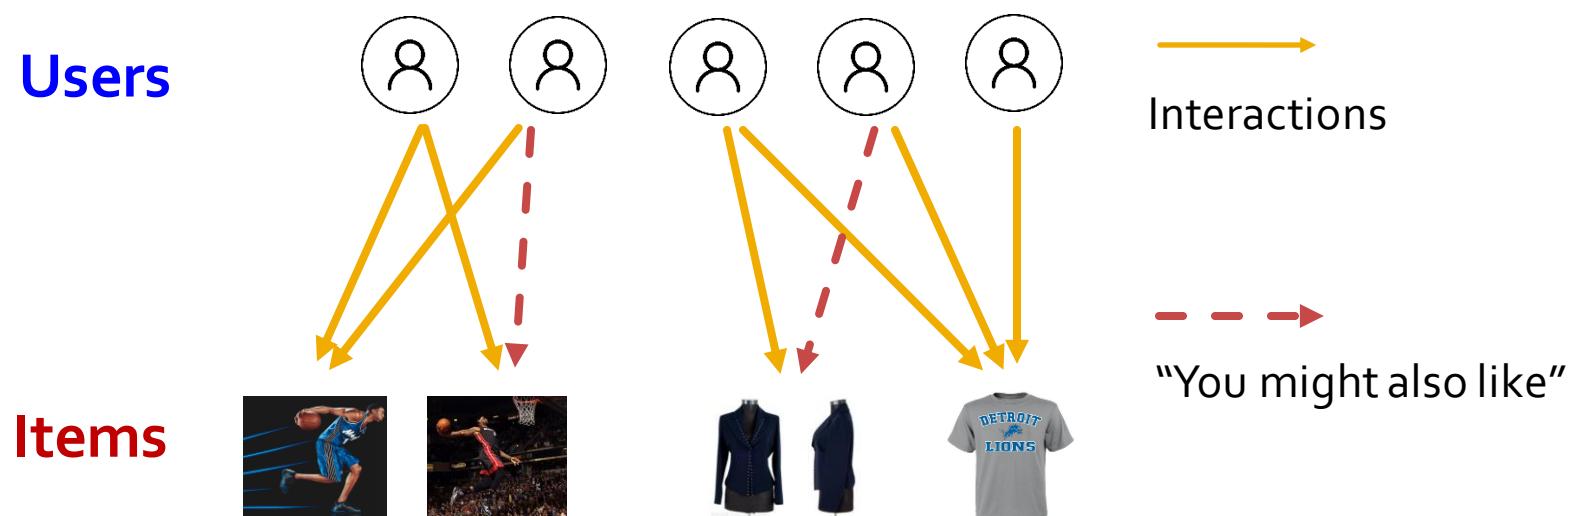
- Key idea: “Spatial graph”
  - Nodes: Amino acids in a protein sequence
  - Edges: Proximity between amino acids (residues)



# Examples of Edge-level ML Tasks

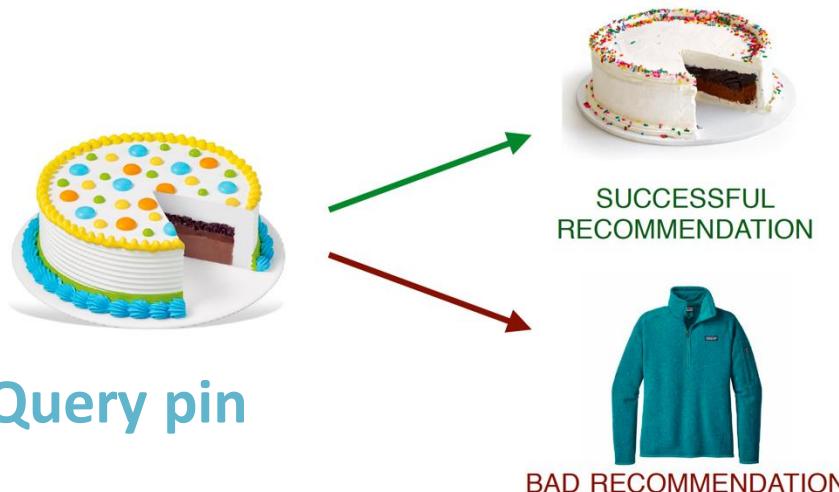
# Example (2): Recommender Systems

- **Users interacts with items**
  - Watch movies, buy merchandise, listen to music
  - **Nodes:** Users and items
  - **Edges:** User-item interactions
- **Goal: Recommend items users might like**



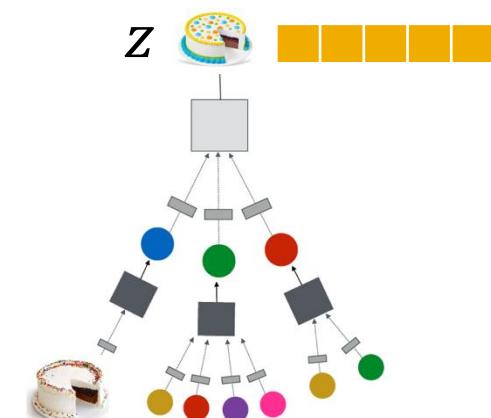
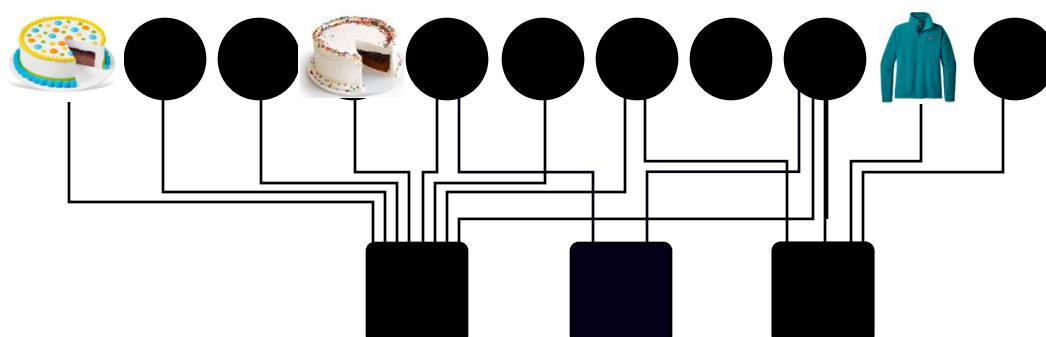
# PinSage: Graph-based Recommender

**Task: Recommend related pins to users**



**Task:** Learn node embeddings  $z_i$  such that  
 $d(z_{cake1}, z_{cake2}) < d(z_{cake1}, z_{sweater})$

**Predict whether two nodes in a graph are related**

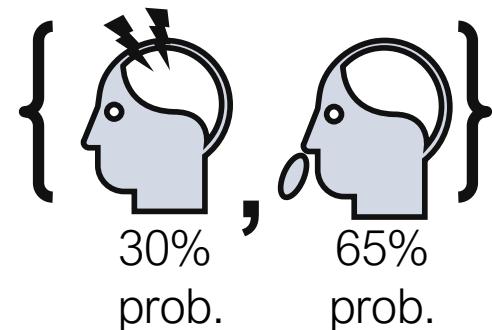
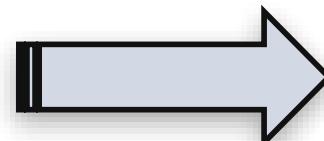


# Example (3): Drug Side Effects

Many patients **take multiple drugs** to treat  
**complex or co-existing diseases:**

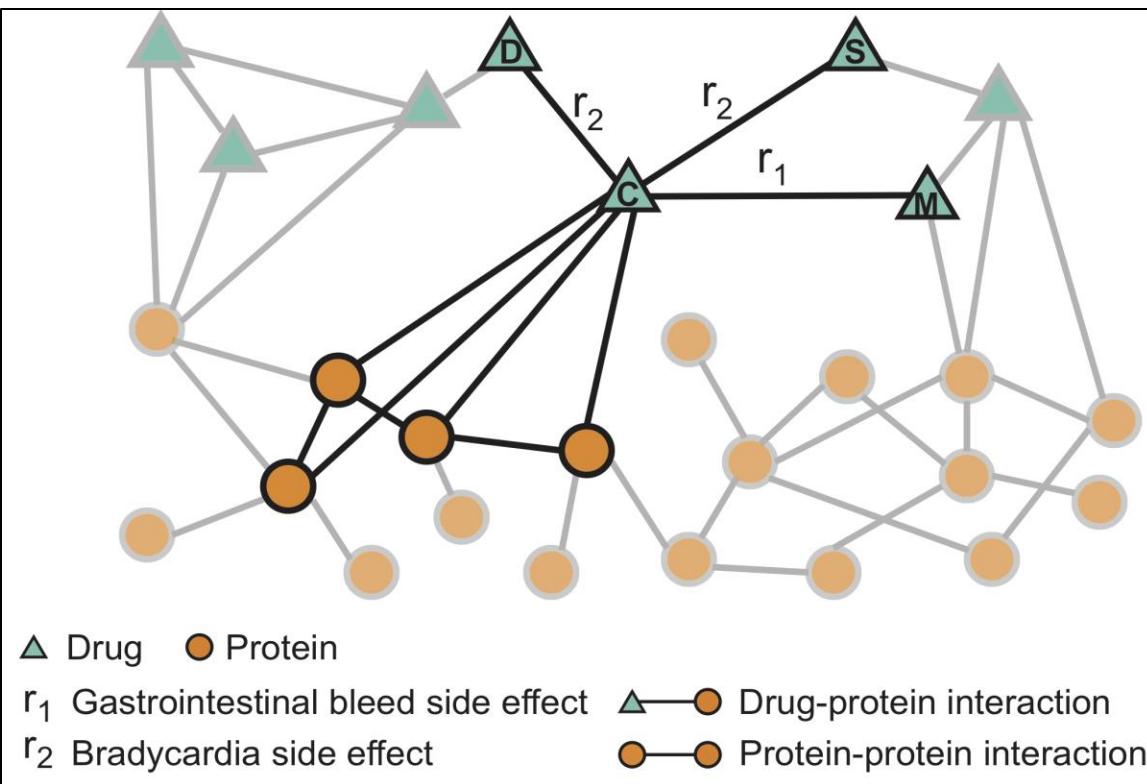
- 46% of people ages 70-79 take more than 5 drugs
- Many patients take more than 20 drugs to treat heart disease, depression, insomnia, etc.

**Task: Given a pair of drugs predict  
adverse side effects**

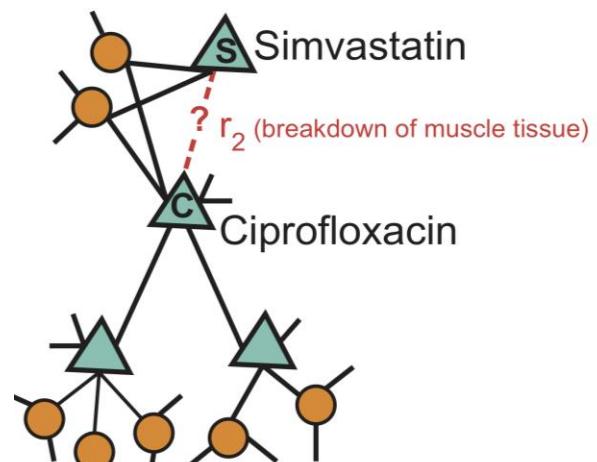


# Biomedical Graph Link Prediction

- **Nodes:** Drugs & Proteins
- **Edges:** Interactions



**Query:** How likely will Simvastatin and Ciprofloxacin, when taken together, break down muscle tissue?



# Results: *De novo* Predictions

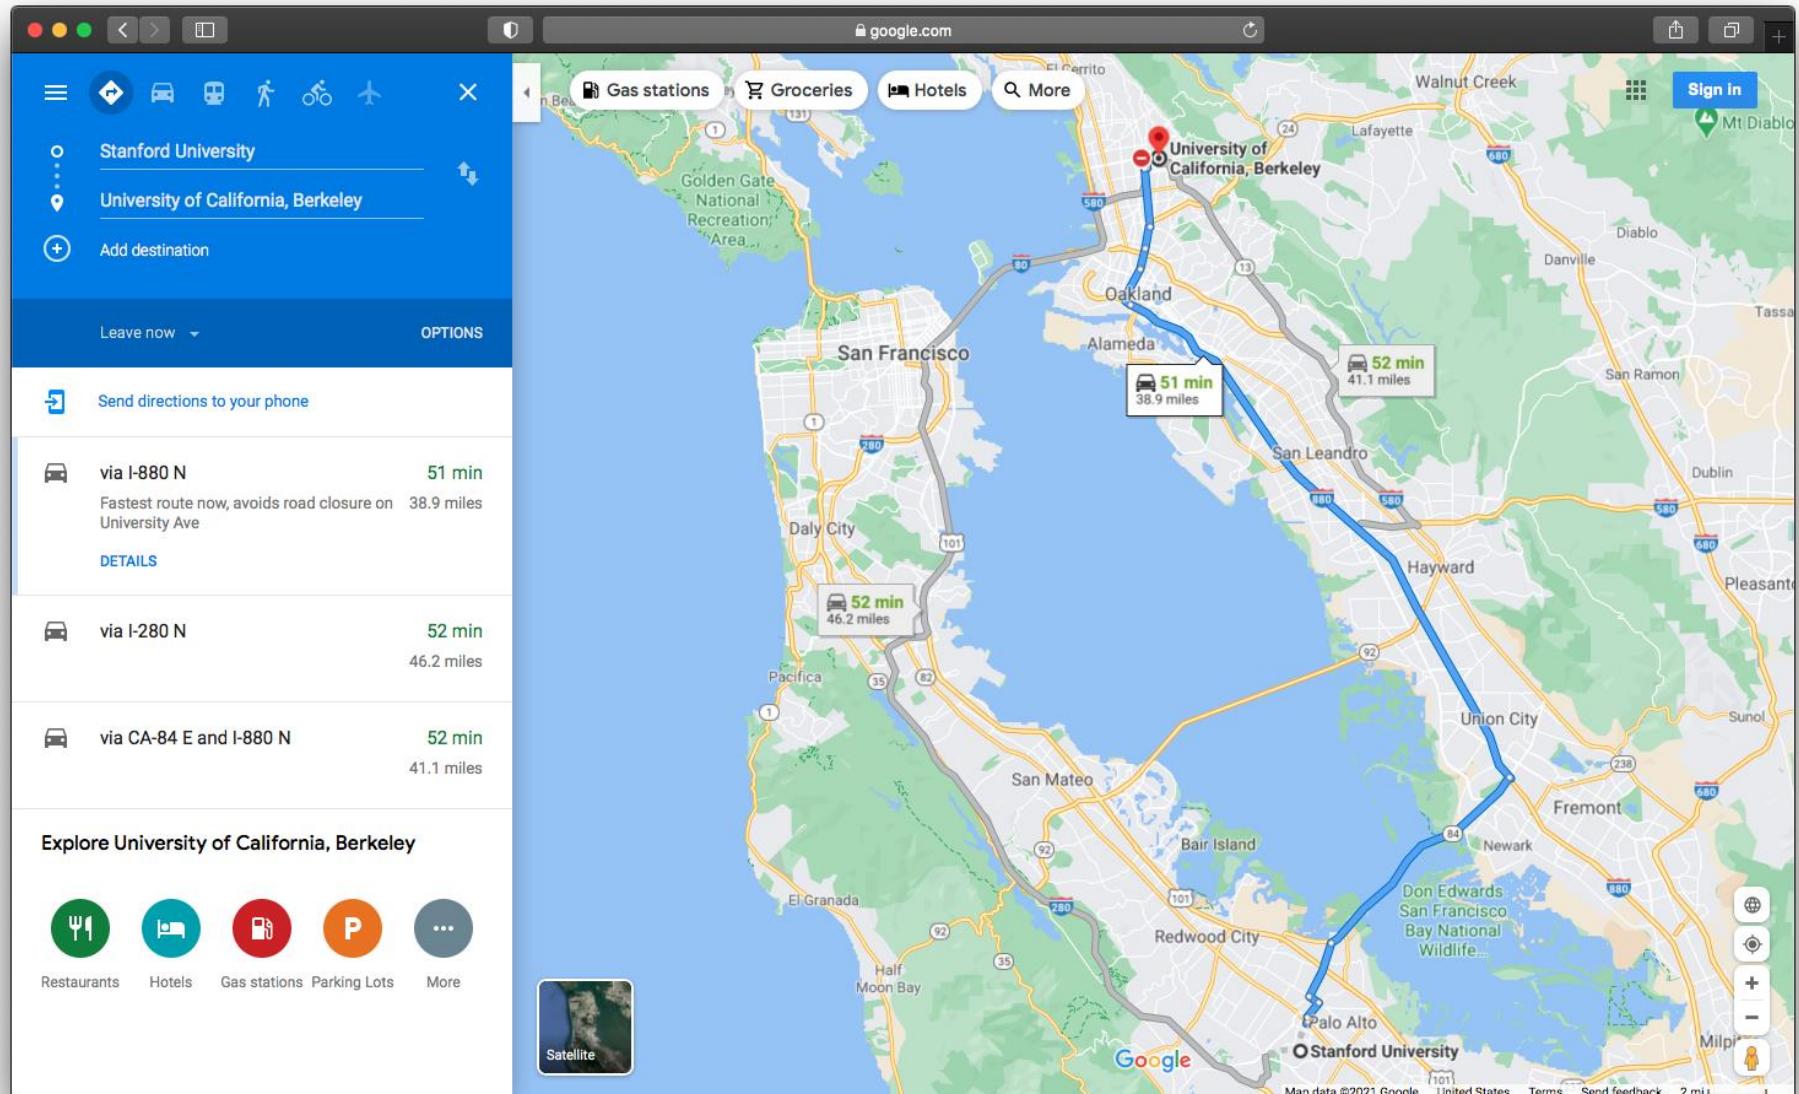
Rank	Drug $c$	Drug $d$	Side effect $r$	Evidence found
1	Pyrimethamine	Aliskiren	Sarcoma	<a href="#">Stage et al. 2015</a>
2	Tigecycline	Bimatoprost	Autonomic neuropathy	
3	Omeprazole	Dacarbazine	Telangiectases	
4	Tolcapone	Pyrimethamine	Breast disorder	<a href="#">Bicker et al. 2017</a>
5	Minoxidil	Paricalcitol	Cluster headache	
6	Omeprazole	Amoxicillin	Renal tubular acidosis	<a href="#">Russo et al. 2016</a>
7	Anagrelide	Azelaic acid	Cerebral thrombosis	
8	Atorvastatin	Amlodipine	Muscle inflammation	<a href="#">Banakh et al. 2017</a>
9	Aliskiren	Tioconazole	Breast inflammation	<a href="#">Parving et al. 2012</a>
10	Estradiol	Nadolol	Endometriosis	

*Case Report*

**Severe Rhabdomyolysis due to Presumed Drug Interactions  
between Atorvastatin with Amlodipine and Ticagrelor**

# Examples of Subgraph-level ML Tasks

# Example (4): Traffic Prediction



# Road Network as a Graph

- **Nodes:** Road segments
- **Edges:** Connectivity between road segments
- **Prediction:** Time of Arrival (ETA)

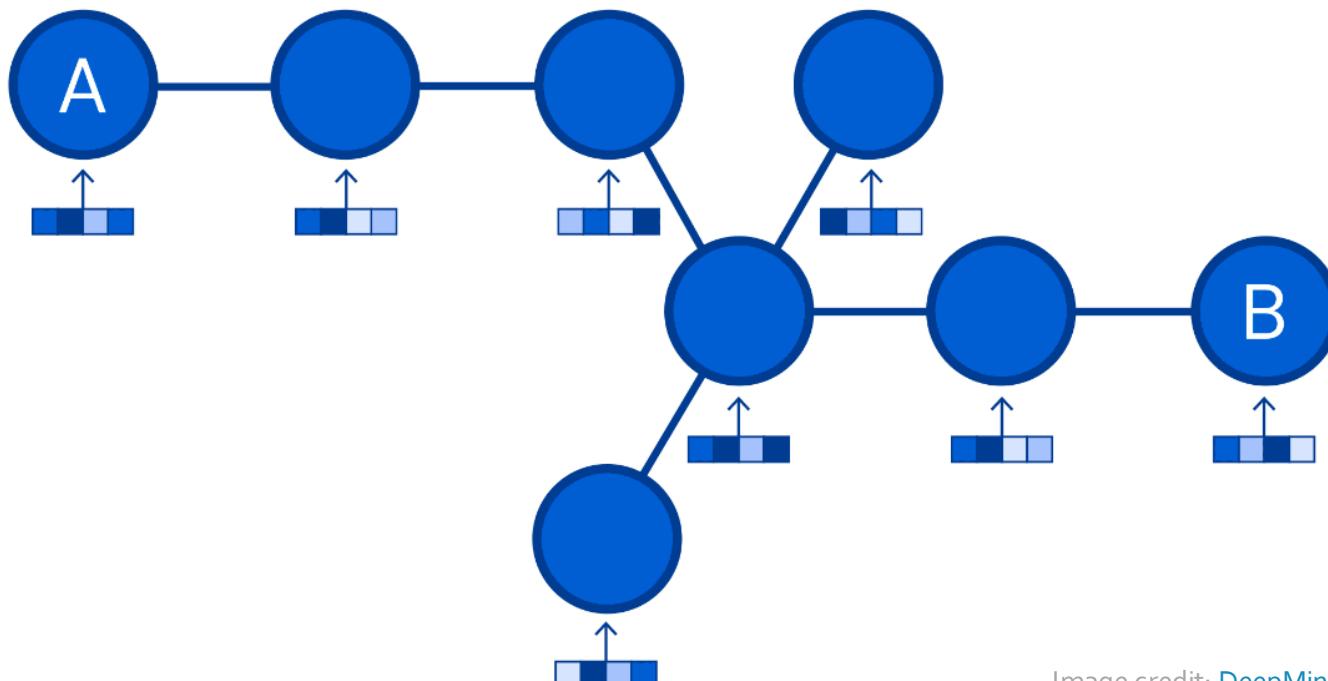
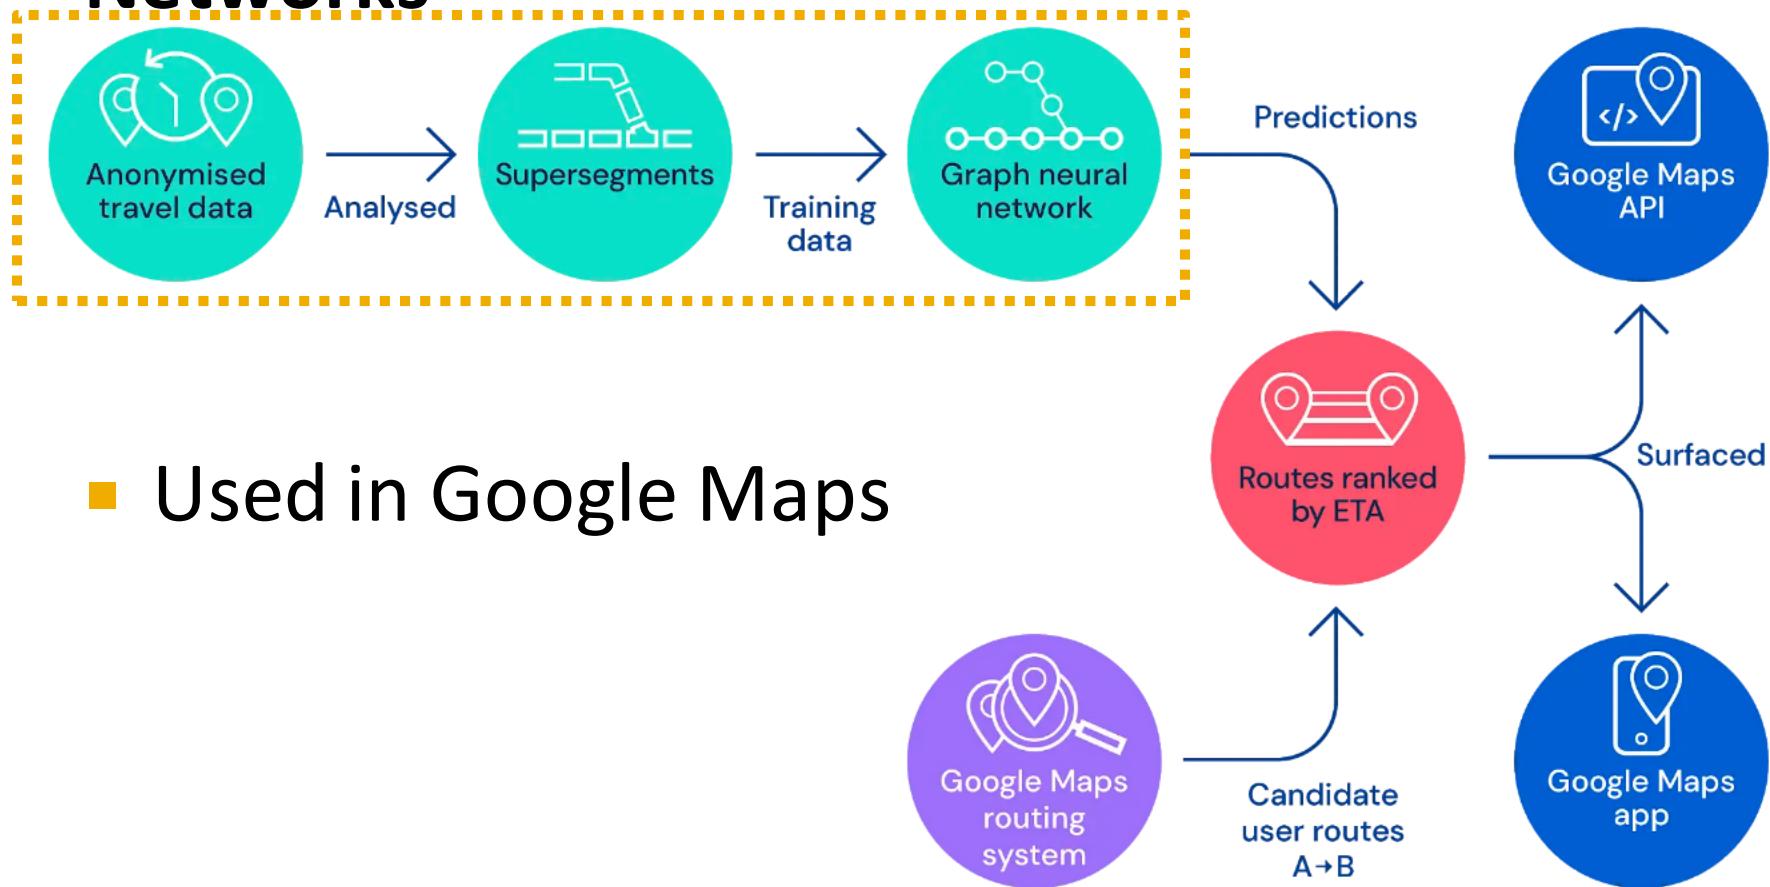


Image credit: [DeepMind](#)

# Traffic Prediction via GNN

## Predicting Time of Arrival with Graph Neural Networks



- Used in Google Maps

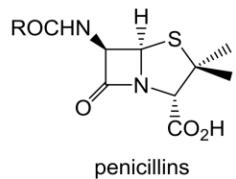
# Examples of Graph-level ML Tasks

# Example (5): Drug Discovery

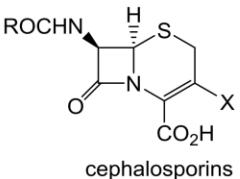
## ■ Antibiotics are small molecular graphs

- **Nodes:** Atoms

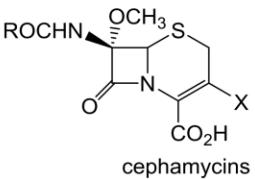
- **Edges:** Chemical bonds



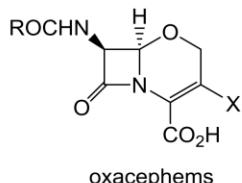
penicillins



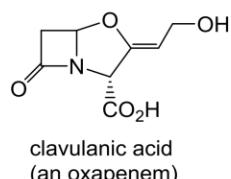
cephalosporins



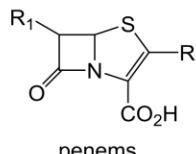
cephamycins



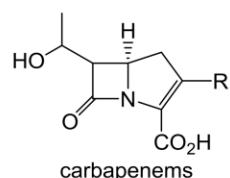
oxacephems



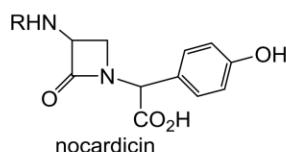
clavulanic acid  
(an oxapenem)



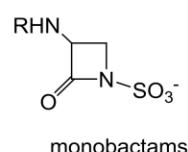
penems



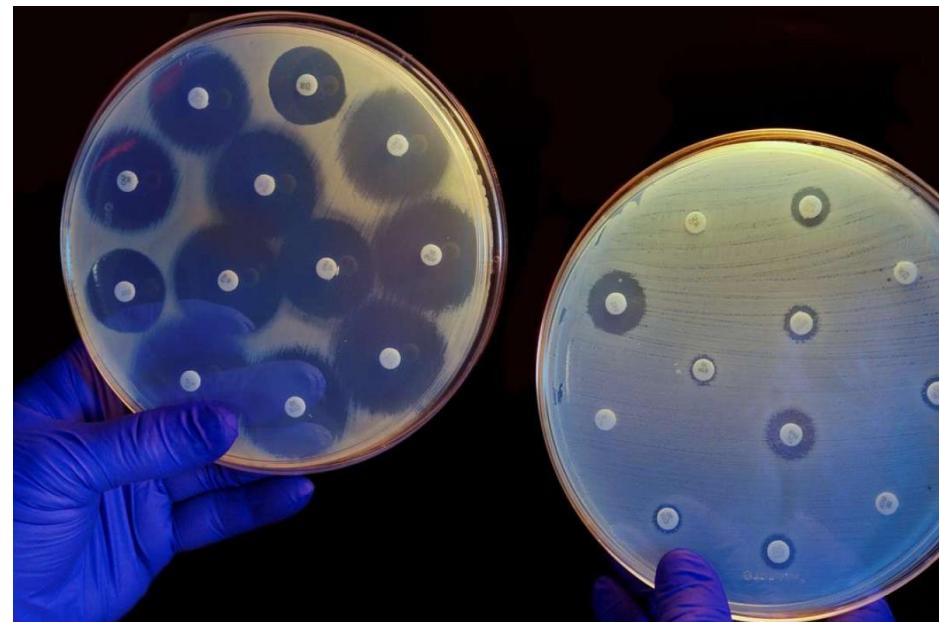
carbapenems



nocardicin



monobactams

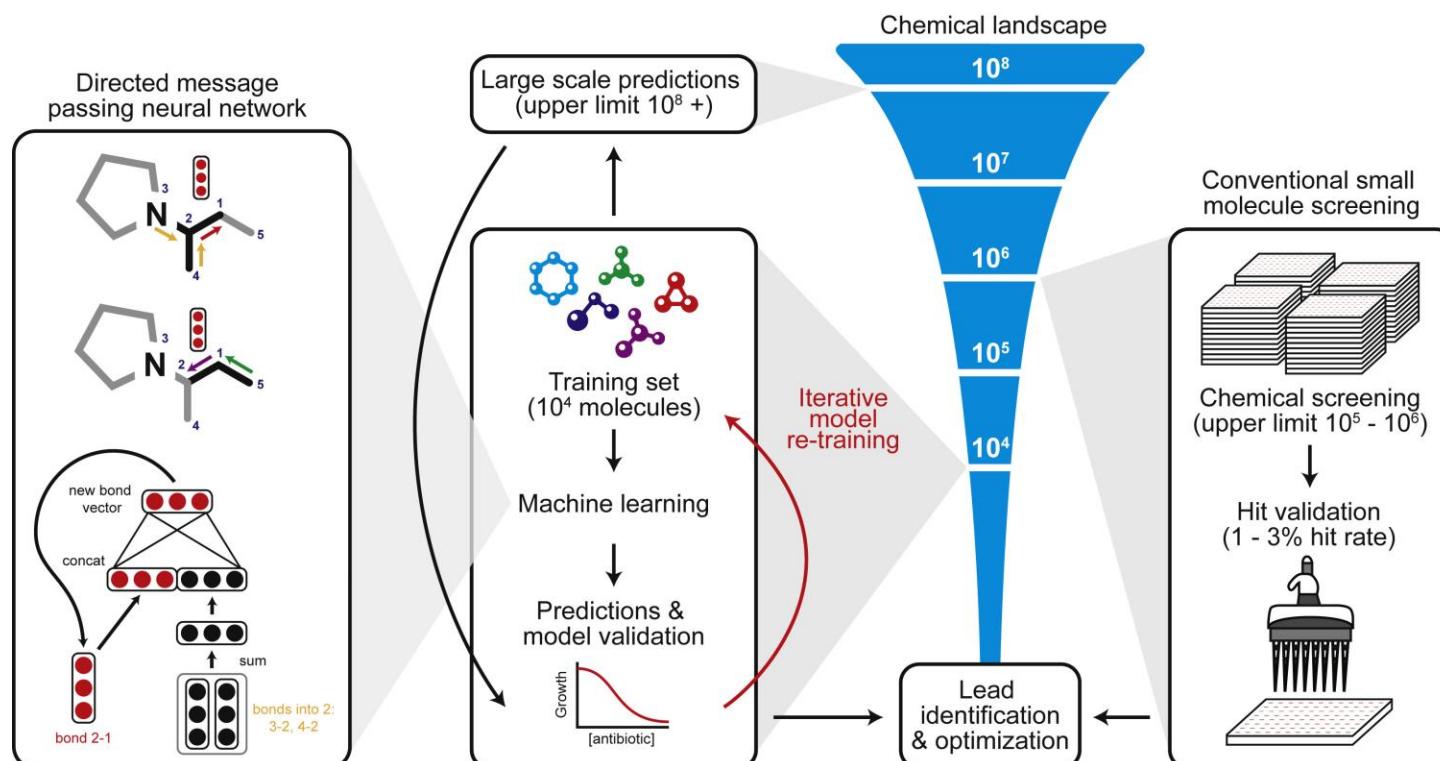


Konaklieva, Monika I. "Molecular targets of  $\beta$ -lactam-based antimicrobials: beyond the usual suspects." *Antibiotics* 3.2 (2014): 128-142.

Image credit: [CNN](#)

# Deep Learning for Antibiotic Discovery

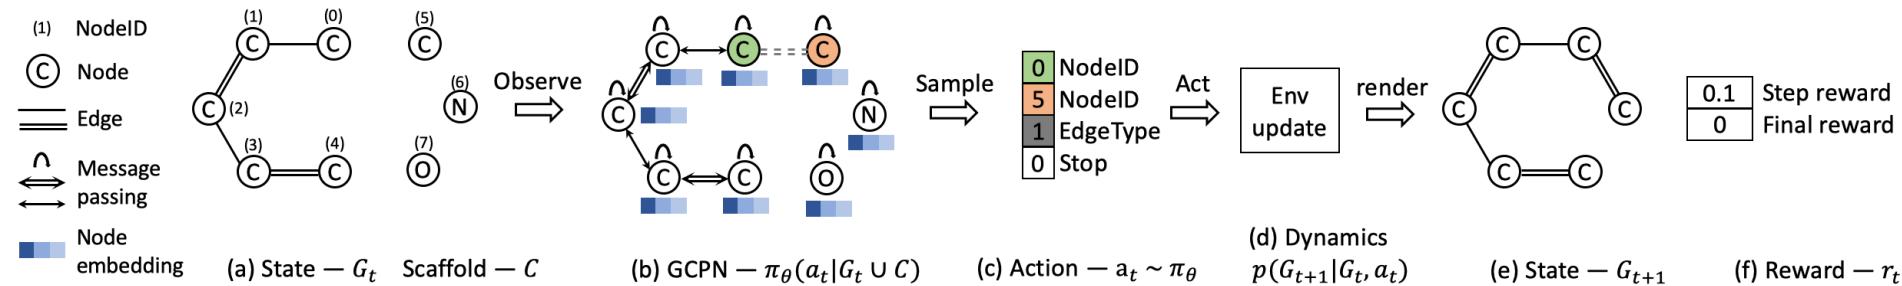
- A Graph Neural Network **graph classification model**
- Predict promising molecules from a pool of candidates



Stokes, Jonathan M., et al. "A deep learning approach to antibiotic discovery." Cell 180.4 (2020): 688-702.

# Molecule Generation / Optimization

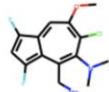
## Graph generation: Generating novel molecules



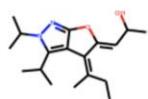
**Use case 1:** Generate novel molecules with high Drug likeness value



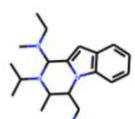
0.948



0.945



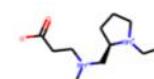
0.944



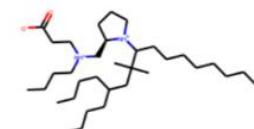
0.941

**Drug likeness**

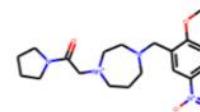
**Use case 2:** Optimize existing molecules to have desirable properties



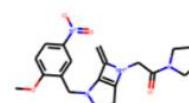
-8.32



-0.71



-5.55

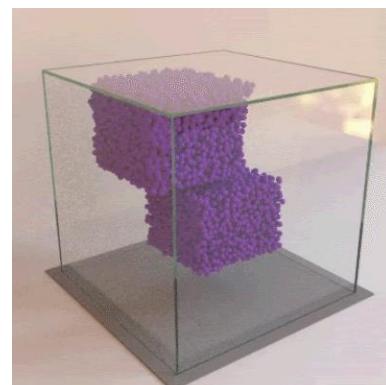
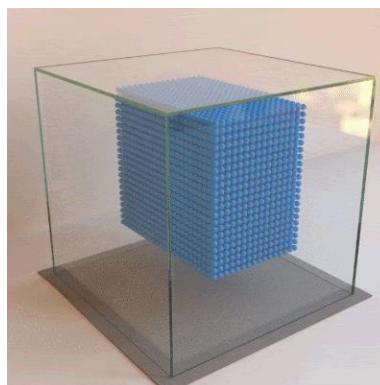


-1.78

# Example (6): Physics Simulation

Physical simulation as a graph:

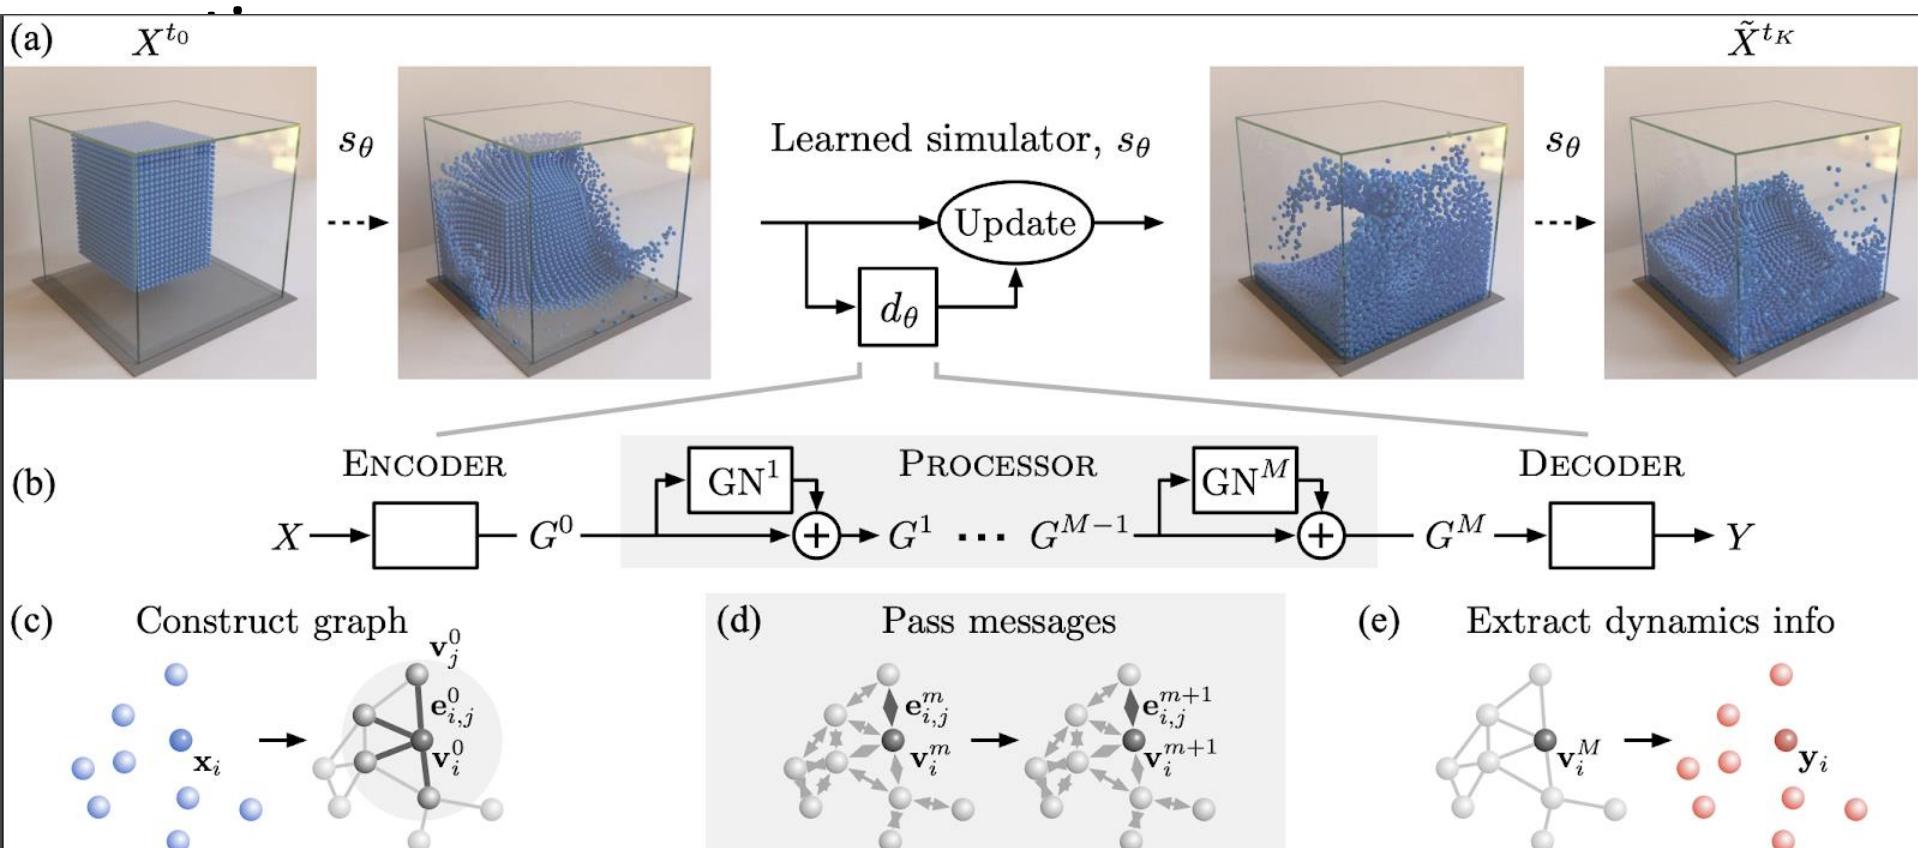
- **Nodes:** Particles
- **Edges:** Interaction between particles



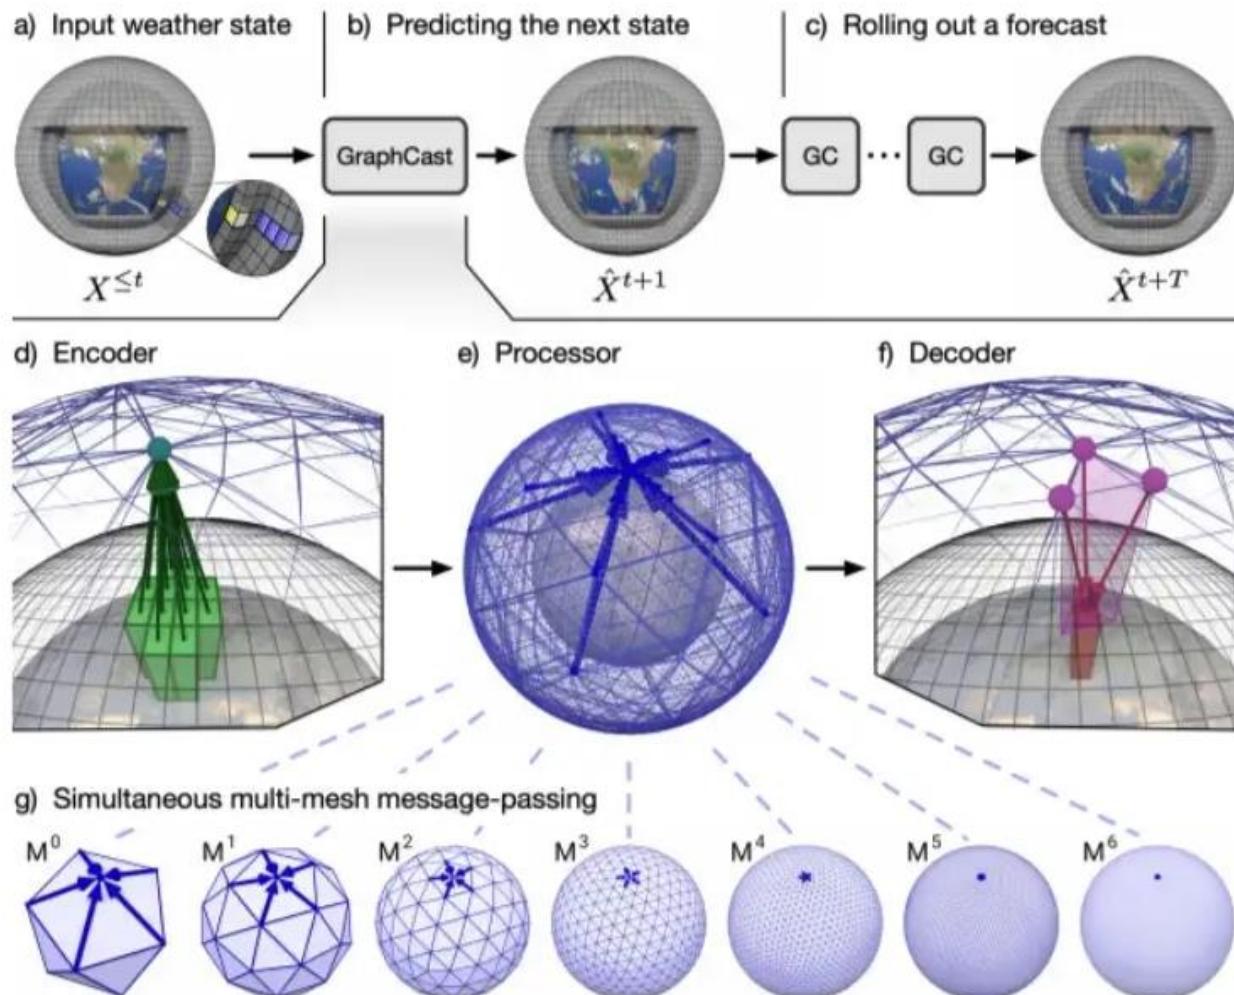
# Simulation Learning Framework

## A graph evolution task:

- **Goal:** Predict how a graph will evolve over time



# Application: DeepMind weather forecasting



<https://medium.com/syncedreview/deepmind-googles-ml-based-graphcast-outperforms-the-world-s-best-medium-range-weather-9d114460aa0c>

# **Stanford CS224W: Choice of Graph Representation**

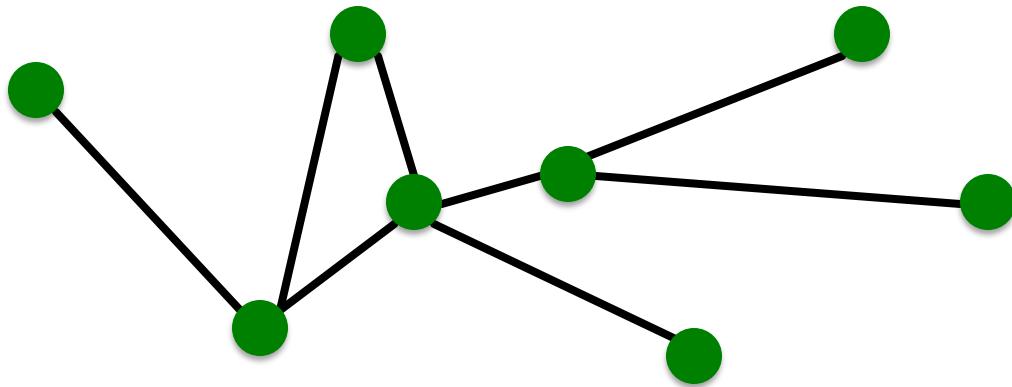
CS224W: Machine Learning with Graphs

Jure Leskovec, Stanford University

<http://cs224w.stanford.edu>

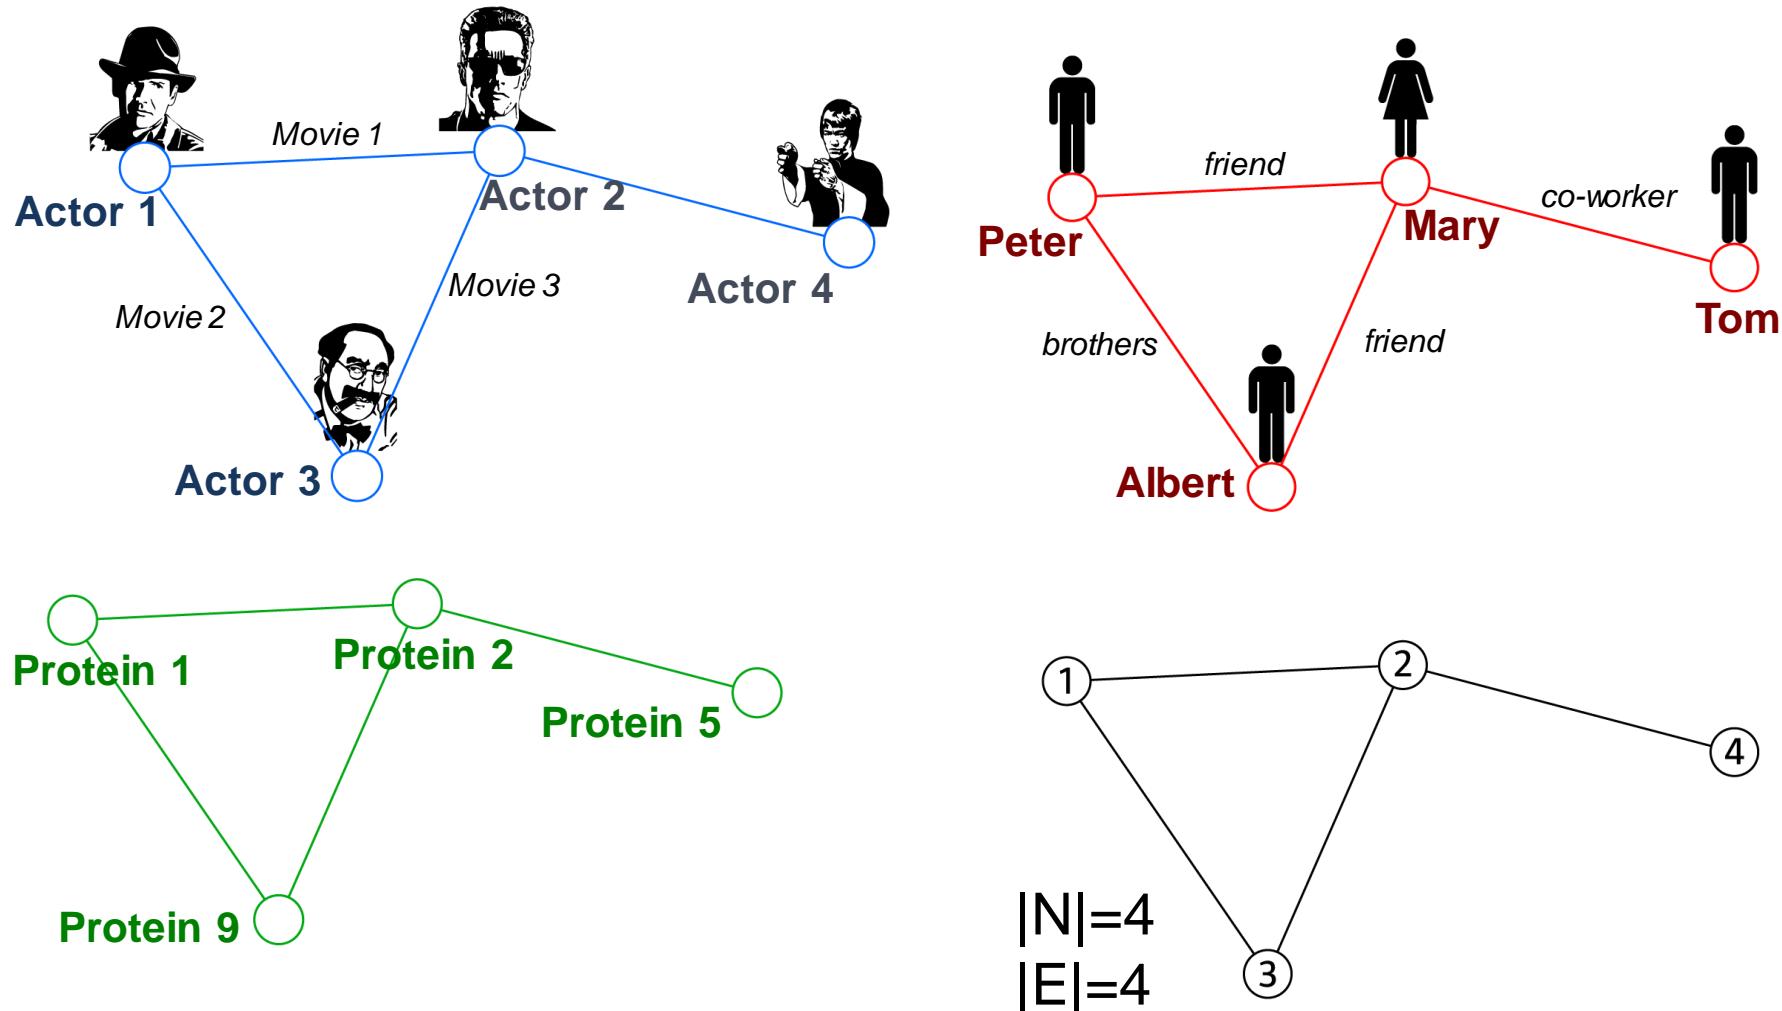


# Components of a Network



- **Objects:** nodes, vertices  $N$
- **Interactions:** links, edges  $E$
- **System:** network, graph  $G(N,E)$

# Graphs: A Common Language



# Choosing a Proper Representation

- If you connect individuals that work with each other, you will explore a **professional network**
- If you connect those that have a sexual relationship, you will be exploring **sexual networks**
- If you connect scientific papers that cite each other, you will be studying the **citation network**
- **If you connect all papers with the same word in the title, what will you be exploring?** It is a network, nevertheless



Image credit: [Euro Scientists](#)

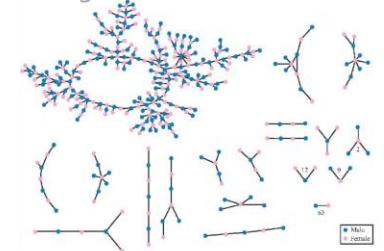
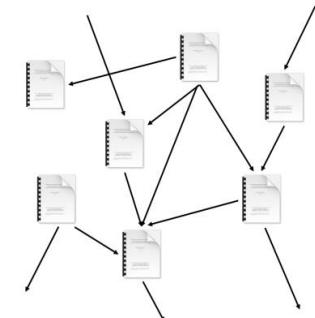


Image credit: [ResearchGate](#)



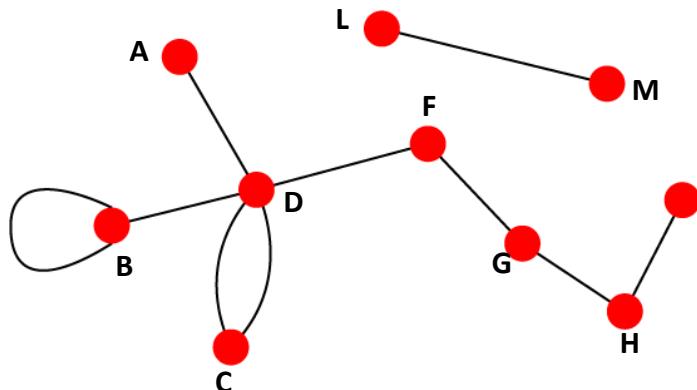
# How do you define a graph?

- **How to build a graph:**
  - What are nodes?
  - What are edges?
- **Choice of the proper network representation of a given domain/problem determines our ability to use networks successfully:**
  - In some cases, there is a unique, unambiguous representation
  - In other cases, the representation is by no means unique
  - The way you assign links will determine the nature of the question you can study

# Directed vs. Undirected Graphs

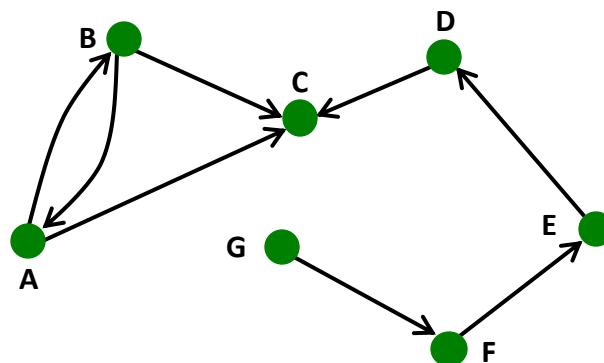
## Undirected

- Links: undirected  
(symmetrical, reciprocal)



## Directed

- Links: directed  
(arcs)



## Examples:

- Collaborations
- Friendship on Facebook

## Examples:

- Phone calls
- Following on Twitter

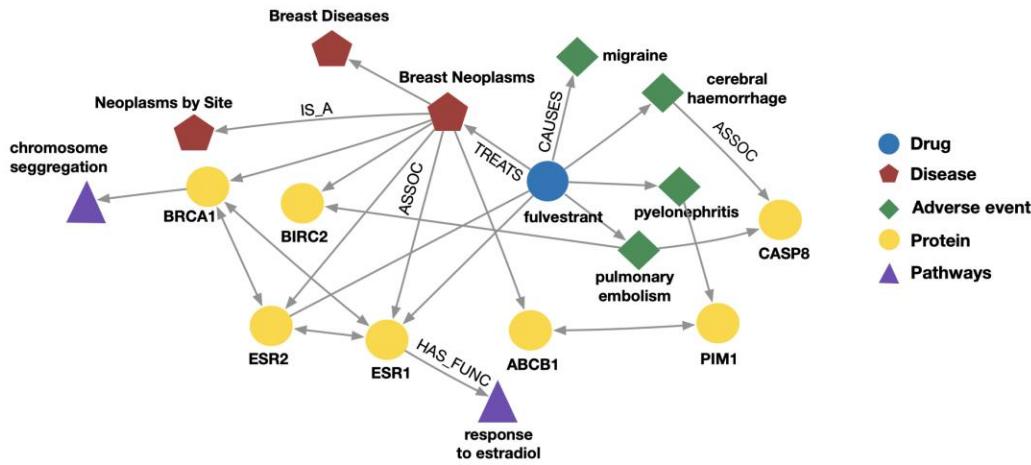
# Heterogeneous Graphs

- A heterogeneous graph is defined as

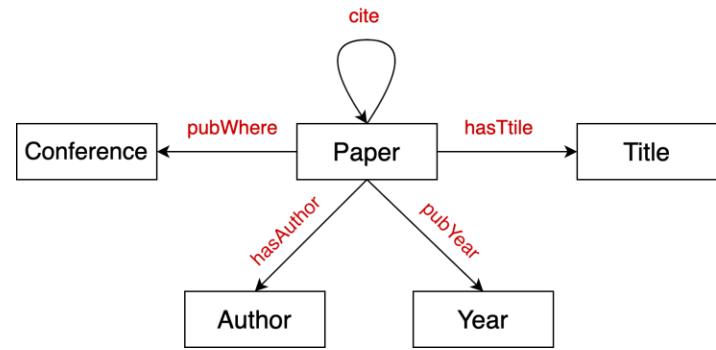
$$G = (V, E, R, T)$$

- Nodes with node types  $v_i \in V$
- Edges with relation types  $(v_i, r, v_j) \in E$
- Node type  $T(v_i)$
- Relation type  $r \in R$

# Many Graphs are Heterogeneous Graphs



- Drug
- Disease
- Adverse event
- Protein
- Pathways



## Biomedical Knowledge Graphs

Example node: Migraine

Example edge: (fulvestrant, Treats, Breast Neoplasms)

Example node type: Protein

Example edge type (relation): Causes

## Academic Graphs

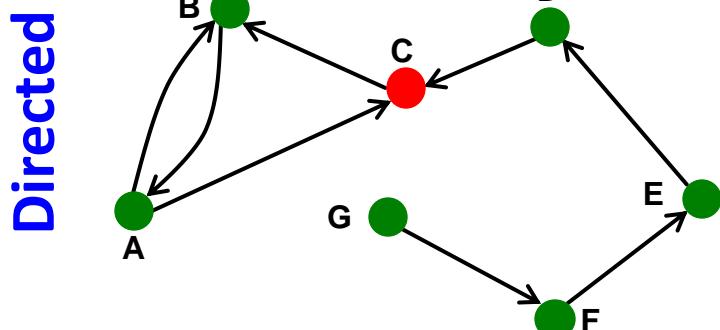
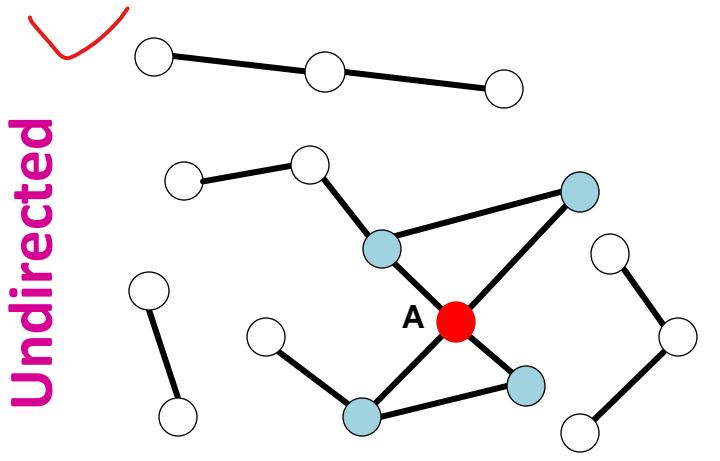
Example node: ICML

Example edge: (GraphSAGE, NeurIPS)

Example node type: Author

Example edge type (relation): pubYear

# Node Degrees



**Source:** Node with  $k^{in} = 0$

**Sink:** Node with  $k^{out} = 0$

**Node degree,  $k_i$ :** the number of edges adjacent to node  $i$

$$k_A = 4$$

**Avg. degree:**  $\bar{k} = \langle k \rangle = \frac{1}{N} \sum_{i=1}^N k_i = \frac{2E}{N}$

In directed networks we define an **in-degree** and **out-degree**. The (total) degree of a node is the sum of in- and out-degrees.

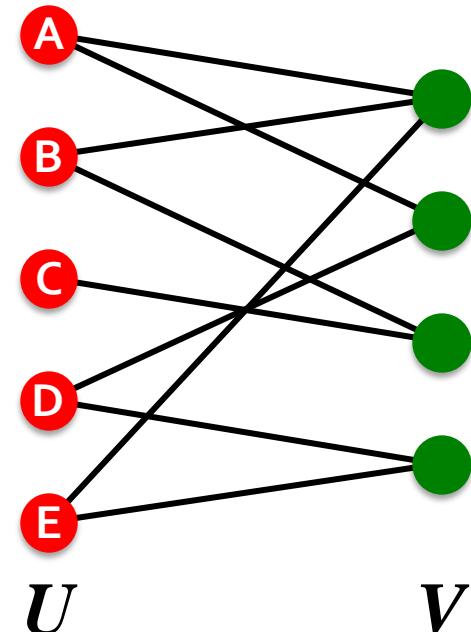
$$k_C^{in} = 2 \quad k_C^{out} = 1 \quad k_C = 3$$

$$\bar{k} = \frac{E}{N}$$

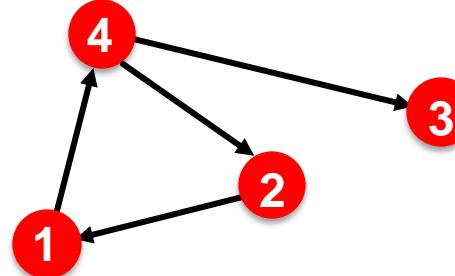
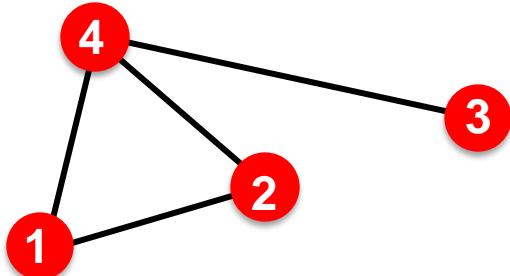
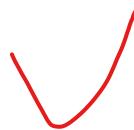
$$\bar{k}^{in} = \bar{k}^{out}$$

# Bipartite Graph

- **Bipartite graph** is a graph whose nodes can be divided into two disjoint sets  $U$  and  $V$  such that every link connects a node in  $U$  to one in  $V$ ; that is,  $U$  and  $V$  are **independent sets**
- **Examples:**
  - Authors-to-Papers (they authored)
  - Actors-to-Movies (they appeared in)
  - Users-to-Movies (they rated)
  - Recipes-to-Ingredients (they contain)
- **“Folded” networks:**
  - Author collaboration networks
  - Movie co-rating networks



# Representing Graphs: Adjacency Matrix



$A_{ij} = 1$  if there is a link from node  $i$  to node  $j$

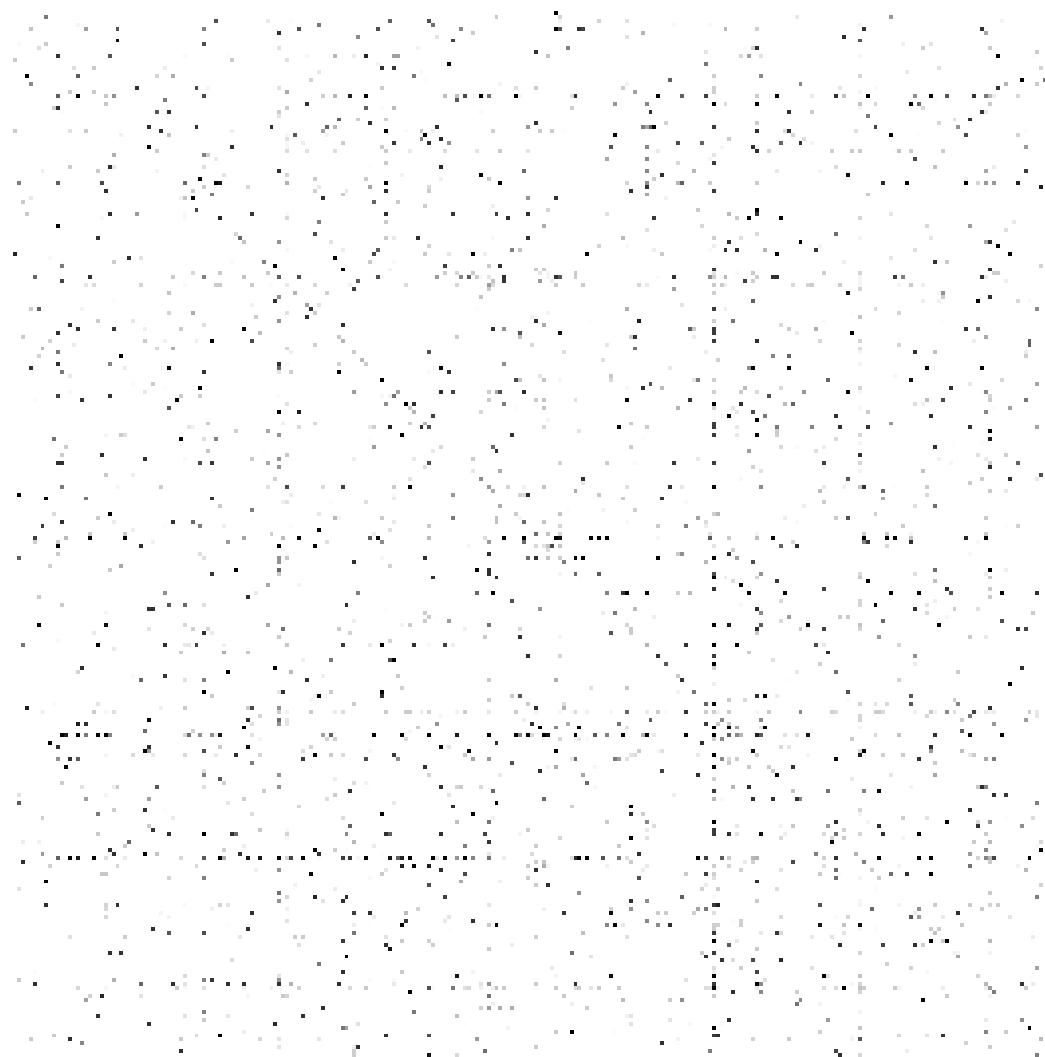
$A_{ij} = 0$  otherwise

$$A = \begin{pmatrix} 0 & 1 & 0 & 1 \\ 1 & 0 & 0 & 1 \\ 0 & 0 & 0 & 1 \\ 1 & 1 & 1 & 0 \end{pmatrix}$$

$$A = \begin{pmatrix} 0 & 0 & 0 & 1 \\ 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 \\ 0 & 1 & 1 & 0 \end{pmatrix}$$

Note that for a directed graph (right) the matrix is not symmetric.

# Adjacency Matrices are Sparse



# Networks are Sparse Graphs

Most real-world networks are **sparse**

$$E \ll E_{\max} \text{ (or } k \ll N-1)$$

NETWORK	NODES	LINKS	DIRECTED/ UNDIRECTED	N	L	$\langle k \rangle$
Internet	Routers	Internet connections	Undirected	192,244	609,066	6.33
WWW	Webpages	Links	Directed	325,729	1,497,134	4.60
Power Grid	Power plants, transformers	Cables	Undirected	4,941	6,594	2.67
Phone Calls	Subscribers	Calls	Directed	36,595	91,826	2.51
Email	Email Addresses	Emails	Directed	57,194	103,731	1.81
Science Collaboration	Scientists	Co-authorship	Undirected	23,133	93,439	8.08
Actor Network	Actors	Co-acting	Undirected	702,388	29,397,908	83.71
Citation Network	Paper	Citations	Directed	449,673	4,689,479	10.43
E. Coli Metabolism	Metabolites	Chemical reactions	Directed	1,039	5,802	5.58
Protein Interactions	Proteins	Binding interactions	Undirected	2,018	2,930	2.90

**Consequence:** Adjacency matrix is filled with zeros!

(Density of the matrix ( $E/N^2$ ): WWW=1.51x10<sup>-5</sup>, MSN IM = 2.27x10<sup>-8</sup>)

# Node and Edge Attributes

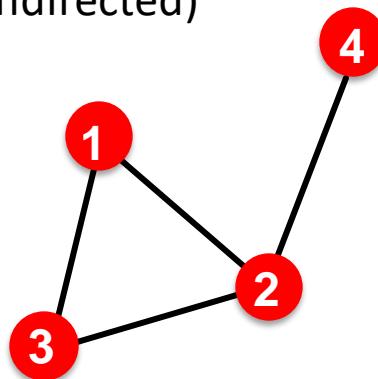
## Possible options:

- Weight (*e.g.*, frequency of communication)
- Ranking (best friend, second best friend...)
- Type (friend, relative, co-worker)
- Sign: Friend vs. Foe, Trust vs. Distrust
- Properties depending on the structure of the rest of the graph: Number of common friends

# More Types of Graphs

## ■ Unweighted

(undirected)



$$A_{ij} = \begin{pmatrix} 0 & 1 & 1 & 0 \\ 1 & 0 & 1 & 1 \\ 1 & 1 & 0 & 0 \\ 0 & 1 & 0 & 0 \end{pmatrix}$$

$$A_{ii} = 0$$

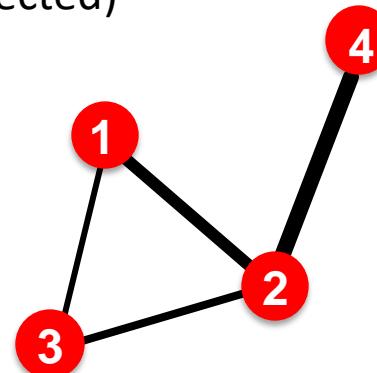
$$A_{ij} = A_{ji}$$

$$E = \frac{1}{2} \sum_{i,j=1}^N A_{ij} \quad \bar{k} = \frac{2E}{N}$$

Examples: Friendship, Hyperlink

## ■ Weighted

(undirected)



$$A_{ij} = \begin{pmatrix} 0 & 2 & 0.5 & 0 \\ 2 & 0 & 1 & 4 \\ 0.5 & 1 & 0 & 0 \\ 0 & 4 & 0 & 0 \end{pmatrix}$$

$$A_{ii} = 0$$

$$A_{ij} = A_{ji}$$

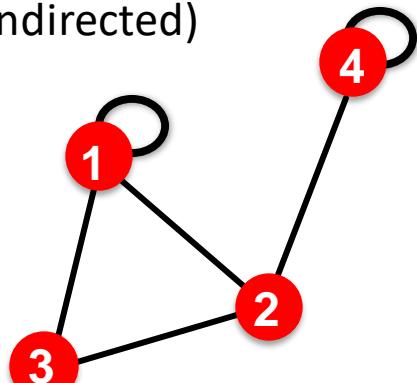
$$E = \frac{1}{2} \sum_{i,j=1}^N \text{nonzero}(A_{ij}) \quad \bar{k} = \frac{2E}{N}$$

Examples: Collaboration, Internet, Roads

# More Types of Graphs

## ■ Self-edges (self-loops)

(undirected)



$$A_{ij} = \begin{pmatrix} 1 & 1 & 1 & 0 \\ 1 & 0 & 1 & 1 \\ 1 & 1 & 0 & 0 \\ 0 & 1 & 0 & 1 \end{pmatrix}$$

$$A_{ii} \neq 0$$

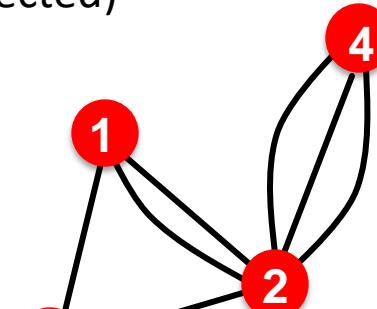
$$A_{ij} = A_{ji}$$

$$E = \frac{1}{2} \sum_{i,j=1, i \neq j}^N A_{ij} + \sum_{i=1}^N A_{ii}$$

Examples: Proteins, Hyperlinks

## ■ Multigraph

(undirected)



$$A_{ij} = \begin{pmatrix} 0 & 2 & 1 & 0 \\ 2 & 0 & 1 & 3 \\ 1 & 1 & 0 & 0 \\ 0 & 3 & 0 & 0 \end{pmatrix}$$

$$A_{ii} = 0$$

$$E = \frac{1}{2} \sum_{i,j=1}^N \text{nonzero}(A_{ij})$$

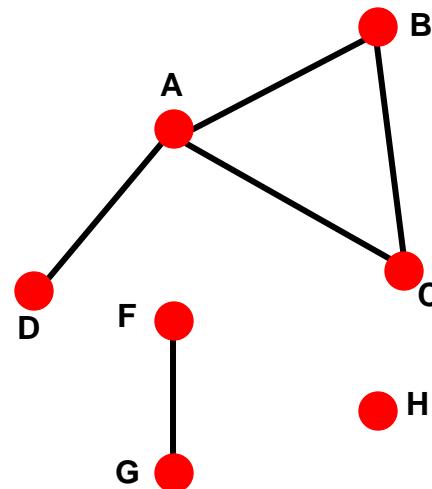
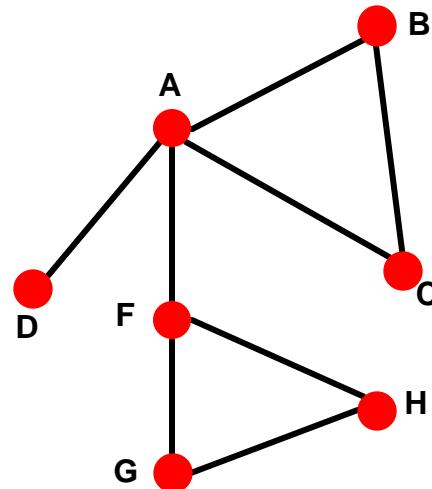
$$A_{ij} = A_{ji}$$

$$\bar{k} = \frac{2E}{N}$$

Examples: Communication, Collaboration

# Connectivity of Undirected Graphs

- **Connected (undirected) graph:**
  - Any two vertices can be joined by a path
- A disconnected graph is made up by two or more connected components



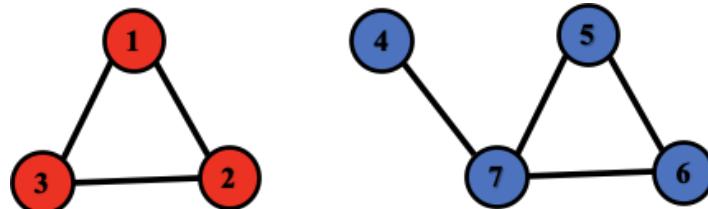
Largest Component:  
**Giant Component**

Isolated node (node H)

# Connectivity: Example

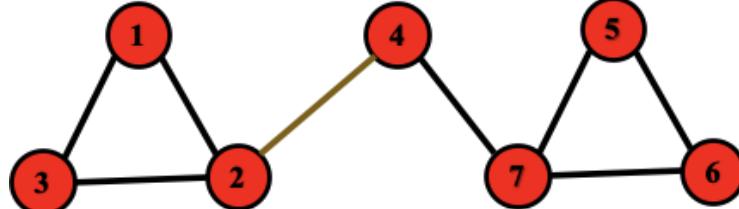
The adjacency matrix of a network with several components can be written in a block-diagonal form, so that nonzero elements are confined to squares, with all other elements being zero:

Disconnected



$$\begin{bmatrix} \begin{matrix} 0 & 1 & 1 \\ 1 & 0 & 1 \\ 1 & 1 & 0 \end{matrix} & \begin{matrix} 0 & 0 & 0 \\ 0 & 0 & 0 \\ 0 & 0 & 0 \end{matrix} \\ \begin{matrix} 0 & 0 & 0 \\ 0 & 0 & 0 \\ 0 & 0 & 0 \end{matrix} & \begin{matrix} 0 & 0 & 0 & 1 \\ 0 & 0 & 1 & 1 \\ 0 & 1 & 0 & 1 \\ 1 & 1 & 1 & 0 \end{matrix} \end{bmatrix}$$

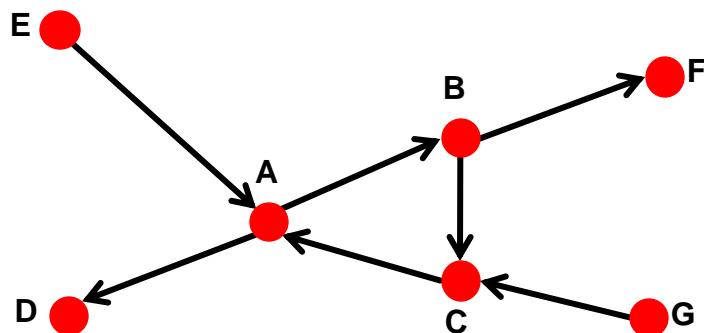
Connected



$$\begin{bmatrix} \begin{matrix} 0 & 1 & 1 \\ 1 & 0 & 1 \\ 1 & 1 & 0 \end{matrix} & \begin{matrix} 0 & 0 & 0 \\ 1 & 0 & 1 \\ 0 & 0 & 0 \end{matrix} \\ \begin{matrix} 0 & 1 & 0 \\ 0 & 0 & 0 \\ 0 & 0 & 0 \end{matrix} & \begin{matrix} 0 & 0 & 0 & 1 \\ 0 & 0 & 1 & 1 \\ 0 & 1 & 0 & 1 \\ 1 & 1 & 1 & 0 \end{matrix} \end{bmatrix}$$

# Connectivity of Directed Graphs

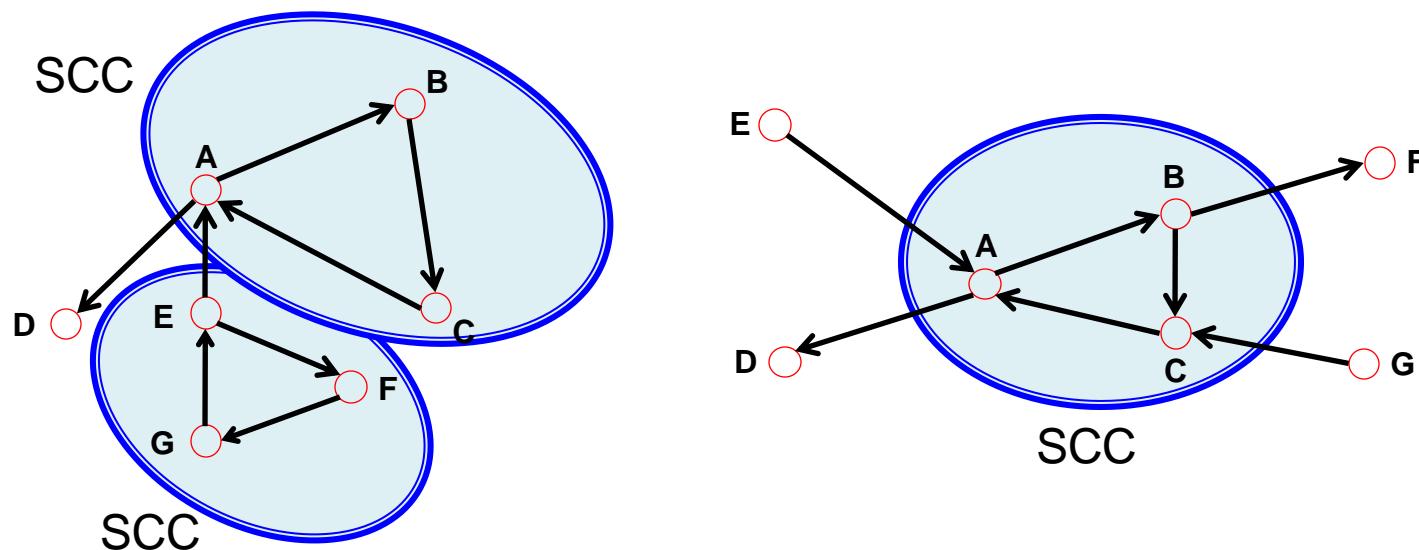
- **Strongly connected directed graph**
  - has a path from each node to every other node and vice versa (e.g., A-B path and B-A path)
- **Weakly connected directed graph**
  - is connected if we disregard the edge directions



Graph on the left is connected but not strongly connected (e.g., there is no way to get from F to G by following the edge directions).

# Connectivity of Directed Graphs

- Strongly connected components (SCCs) can be identified, but not every node is part of a nontrivial strongly connected component.



**In-component:** nodes that can reach the SCC,

**Out-component:** nodes that can be reached from the SCC.

# Summary

- **Machine learning with Graphs**
  - Applications and use cases
- **Different types of tasks:**
  - Node level
  - Edge level
  - Graph level
- **Choice of a graph representation:**
  - Directed, undirected, bipartite, weighted, adjacency matrix

# Stanford CS224W: Course Logistics

CS224W: Machine Learning with Graphs

Jure Leskovec, Stanford University

<http://cs224w.stanford.edu>



# Prerequisites

- **The course is self-contained.**
- **No single topic is too hard by itself.**
- **But we will cover and touch upon many topics and this is what makes the course hard.**
  - **Good background in:**
    - Machine Learning
    - Algorithms and graph theory
    - Probability and statistics
  - **Programming:**
    - You should be able to write non-trivial programs (in Python)
    - Familiarity with PyTorch is a plus

# Graph Machine Learning Tools

- We use [PyG \(PyTorch Geometric\)](#):  PyG
  - The ultimate library for Graph Neural Networks
- We further recommend:
  - [GraphGym](#): Platform for designing Graph Neural Networks.
    - Modularized GNN implementation, simple hyperparameter tuning, flexible user customization
  - Both platforms are very helpful for the course project (save your time & provide advanced GNN functionalities)
- Other network analytics tools: SNAP.PY, NetworkX

# CS224W Course Logistics

- The class meets Tue and Thu 3:00-4:20pm  
**Pacific Time *in person***
  - Videos of the lectures will be recorded and posted on Canvas
- **Structure of lectures:**
  - 70-80 minutes of a lecture
    - During this time you can ask questions
  - 10 minutes of a live Q&A/discussion session at the end of the lecture

# Logistics: Teaching Staff

Instructor



Jure Leskovec

Course Assistants



Hamed Nilforoshan  
Head CA



Serina Chang



Aman Bansal

Guest Instructor



Jiaxuan You



Mohammad Aljubran



Lun Yu (Tina) Li



Paridhi Maheshwari



Sharmila Nangi



Feiyang (Kathy) Yu

# Logistics: Website

- <http://cs224w.stanford.edu>
  - Slides posted before the class
- **Readings:**
  - [Graph Representation Learning Book](#) by Will Hamilton
  - Research papers
- **Optional readings:**
  - Papers and pointers to additional literature
  - **This will be very useful for course projects**

# Logistics: Communication

- **Ed Discussion:**
  - Access via link on Canvas
  - **Please participate and help each other!**
    - Don't post code, annotate your questions, search for answers before you ask
  - We will post course announcements to Ed (make sure you check it regularly)
- **Please don't communicate with prof/TAs via personal emails, but always use:**
  - [cs224w-win2223-staff@lists.stanford.edu](mailto:cs224w-win2223-staff@lists.stanford.edu)

# Logistics: Office Hours

- **OHs will be virtual**
  - We will have OHs every day, starting from 2<sup>nd</sup> week of the course
  - See <http://web.stanford.edu/class/cs224w/oh.html> for Zoom links and link to QueueStatus
  - Schedule to be announced by end of week

# Work for Course: Grading

- **Final grade will be composed of:**
  - **Homework: 25%**
    - 3 written homeworks, each worth 8.3%
  - **Coding assignments: 20%**
    - 5 coding assignments using Google Colab, each worth 4%
  - **Exam: 35%**
  - **Course project: 20%**
    - Proposal: 20%; Final report: 70%; Poster: 10%
  - **Extra credit: Ed participation, PyG/GraphGym code contribution**
    - Used if you are on the boundary between grades

# Work for Course: Submitting

- **How to submit?**
  - **Upload via Gradescope**
    - You will be automatically registered to Gradescope once you officially enroll in CS224W
    - Homeworks, Colabs (numerical answers), and project deliverables are submitted on Gradescope
- **Total of 2 Late Periods (LP) per student**
  - Max 1 LP per assignment (no LP for the final report)
    - LP gives **4 extra days**: assignments usually due on Thursday (11:59pm) → with LP, it is due the following Monday (11:59pm)

# Work for Course: HWs, Colabs

- **Homeworks (25%, n=3)**
  - Written assignments take longer and take time (~10-20h) – start early!
    - A combination of theory, algorithm design, and math
- **Colabs (20%, n=5)**
  - We have more Colabs but they are shorter (~3-5h); Colab 0 is not graded.
    - Get hands-on experience coding and training GNNs; good preparation for final projects and industry

# Work for Course: Exam

- **Single exam: Thursday, March 7 (no class)**
  - Take-home, open-book, timed
    - Administered via Gradescope
    - Released at 10am PT on Thursday March 7, available until 10am PT the following day
    - Once you open it, you will have 100 minutes to complete the exam
  - Content
    - Will have written questions (similar to Homework), will possibly have a coding section (similar to Colabs)
    - More details to come!

# Work for Course: Project

- **Details will be posted soon:**
  - Focus is on real-world applications of GNNs
- **Logistics**
  - **Groups of up to 3 students**
  - Groups of 1 or 2 are allowed; the team size will be taken under consideration when evaluating the scope of the project. But 3 person teams can be more efficient.
  - **Google Cloud credits**
    - We will provide \$50 in Google Cloud credits to each student
    - You can also get \$300 with Google Free Trial (<https://cloud.google.com/free/docs/gcp-free-tier>)
- **Read:** <http://cs224w.stanford.edu/info.html>

# Course Schedule

Assignment	Due on (11:59pm PT)
Colab 0	Not graded
Colab 1	Thu, 1/26 (week 3)
Homework 1	Thu, 2/2 (week 4)
Project Proposal	Tue, 2/7 (week 5)
Colab 2	Thu, 2/9 (week 5)
Homework 2	Thu, 2/16 (week 6)
Colab 3	Tue, 2/23 (week 7)
Homework 3	Thu, 3/2 (week 8)
EXAM	Thu, 3/7 (week 9)
Colab 4	Thu, 3/9 (week 9)
Colab 5	Thu, 3/14 (week 10)
Project Report	Thu, 3/21 (No Late Periods!)

# Honor Code

Make sure you read  
and understand it!

- We strictly enforce the Stanford Honor Code
  - Violations of the Honor Code include:
    - Copying or allowing another to copy from one's own paper
    - Unpermitted collaboration
    - Plagiarism
    - Giving or receiving unpermitted aid on a take-home examination
    - Representing as one's own work the work of another
    - Giving or receiving aid on an assignment under circumstances in which a reasonable person should have known that such aid was not permitted
  - The standard sanction for a first offense includes a one-quarter suspension and 40 hours of community service.

# Course Logistics: Q&A

Two ways to ask questions during lecture:

- **In-person (encouraged)**
- **On Ed:**
  - At the beginning of class, we will open a new discussion thread dedicated to this lecture
  - When to ask on Ed?
    - If you are watching the **livestream** remotely
    - If you have a minor clarifying question
    - If we run out of time to get to your question live
    - **Otherwise, try raising your hand first!**

# Course Logistics: Colab 0

- **Colabs 0 and 1 will be released on our course website at 3pm Thursday (1/12)**
- **Colab 0:**
  - Does not need to be handed-in
- **Colab 1:**
  - Due on Thursday 10/07 (2 weeks from today)
  - Submit written answers and code on Gradescope
  - Will cover material from Lectures 1-4, but you can get started right away!